

STARK-Core-v1

STARK Core Language Specification Overview

Introduction

This document provides an overview of the complete STARK core language specification. The numbered source documents 00–07 plus `CORE-V1-ABSTRACT-MACHINE.md` in `docs/spec/` are normative for Core v1. The concise summary and generated combined artifacts are non-normative views, and compiler-governance ledgers under `docs/compiler/semantic-freeze/` are non-normative. The core language defines the general-purpose language surface (lexing, syntax, types, semantics, memory, modules, and standard library). Non-core extensions are defined separately.

Design Philosophy

Core Principles

1. **Memory Safety:** Prevent common memory errors through ownership and borrowing
2. **Performance:** Zero-cost abstractions and predictable execution
3. **Clarity:** Simple, explicit syntax and semantics
4. **Pragmatism:** A minimal, implementable Core v1
5. **Interoperability:** Clear boundaries for future extensions

Language Goals

- A safe, compiled, general-purpose language core
- Compile-time guarantees for memory and type safety
- Simple, predictable semantics suitable for tooling
- Clear extension points for domain-specific features

Specification Structure

1. Lexical Grammar (01-Lexical-Grammar.md)

Defines how source code is tokenized: - **Keywords**: Control flow, declarations, types, operators - **Identifiers**: Variables, functions, types (snake_case/PascalCase) - **Literals**: Integers, floats, strings, characters, booleans - **Operators**: Arithmetic, comparison, logical, bitwise, assignment - **Comments**: Single-line (//) and multi-line (/* */) - **Whitespace**: Space, tab, newline handling

2. Syntax Grammar (02-Syntax-Grammar.md)

Defines the concrete syntax using EBNF: - **Program Structure**: Items (functions, structs, enums, traits) - **Expressions**: Precedence, associativity - **Statements**: Variable declarations, control flow, returns - **Type Syntax**: Primitives, composites, references, functions - **Pattern Matching**: Destructuring and exhaustiveness

3. Type System (03-Type-System.md)

Comprehensive type system with safety guarantees: - **Primitive Types**: Integers, floats, booleans, characters, strings - **Composite Types**: Arrays, tuples, structs, enums - **Reference Types**: Immutable and mutable references - **Ownership Model**: Move semantics, borrowing rules - **Type Inference**: Local type inference (function signatures are fully explicit) - **Trait System**: Interfaces and generic constraints

4. Semantic Analysis (04-Semantic-Analysis.md)

Rules for meaningful program validation: - **Symbol Resolution**: Scoping, shadowing, name lookup - **Type Checking**: Assignment compatibility, function calls - **Ownership Analysis**: Move tracking, borrow checking - **Control Flow**: Reachability, return path analysis - **Pattern Exhaustiveness**: Match completeness checking - **Error Reporting**: Comprehensive diagnostics with suggestions

5. Memory Model (05-Memory-Model.md)

Memory-safety guarantees and their authority boundaries: - **Ownership Rules**: Single ownership, automatic cleanup - **Move Semantics**: Explicit ownership transfer - **Borrowing System**: Immutable and mutable references - **Lifetime Tracking**: Reference validity guarantees - **Allocation Independence**: Safety rules do not promise a physical allocation strategy - **Drop System**: Deterministic resource cleanup

Abstract Machine (CORE-V1-ABSTRACT-MACHINE.md)

Defines backend-independent execution: - **Values and Places**: Abstract objects, storage, owners, and projections - **Evaluation**: Exact order, exactly-once evaluation, and normal control transfer - **Moves and Destruction**: Replacement, partial moves,

temporaries, loops, and collections - **References**: Identity-preserving projection, returned references, and slice views - **Traps and Observations**: Abort boundary and the differential execution comparator

6. Standard Library ([06-Standard-Library.md](#))

Essential types and functions for practical programming: - **Core Types**: Option, Result, Box, Vec, HashMap - **String Handling**: Unicode-aware string operations - **Collections**: Dynamic arrays, hash tables, sets - **IO Operations**: File handling, console output - **Math Functions**: Arithmetic, trigonometric, random numbers - **Error Handling**: Structured error types and propagation

7. Modules and Packages ([07-Modules-and-Packages.md](#))

Defines module structure, visibility, and import resolution: - **Modules**: File and directory layout - **Visibility**: pub/priv rules - **Imports**: use paths, trees, and aliasing - **Packages**: Manifest and dependency resolution

8. Non-Core Extensions ([../extensions/](#))

Non-core language extensions live outside Core v1 and are optional. The tensor & model type system is specified in docs/extensions/Tensor-Model-Types.md; the remaining AI/ML surface is sketched in docs/extensions/AI-Extensions.md.

Core Language Features

Variables and Mutability

```
let x = 42;           // Immutable by default
let mut y = 10;       // Explicitly mutable
const MAX_SIZE: Int32 = 1000; // Compile-time constant
```

Functions

```
fn add(a: Int32, b: Int32) -> Int32 {
    a + b
}

fn greet(name: &str) {
    println(name);
}
```

Ownership and Borrowing

```
fn consume(s: String) {
    // s is owned here
}
```

```
fn borrow(s: &String) -> UInt64 {
    s.len()
}
```

```
fn mutate(s: &mut String) {
    s.push('!');
}
```

This specification provides a focused foundation for implementing a safe, performant, general-purpose language core.

Conformance

A conforming Core v1 implementation **MUST** follow the requirements in this document. Any deviations or extensions **MUST** be explicitly documented by the implementation.

STARK Lexical Grammar Specification

Overview

This document defines the lexical structure of STARK - how source code is broken down into tokens.

Token Categories

1. Keywords

```
// Control Flow
if, else, match
for, while, loop, break, continue, return
```

```
// Declarations
fn, struct, enum, trait, impl, let, mut, const, type, use, mod
```

```
// Types
Int8, Int16, Int32, Int64, UInt8, UInt16, UInt32, UInt64
Float32, Float64, Bool, String, Char, Unit, str
```

```
// Visibility
pub, priv
```

```
// Module Paths and Self Type
```

self, Self, super, crate

// Operators
in, as

// Literals
true, false

2. Identifiers

IDENTIFIER := [a-zA-Z_] [a-zA-Z0-9_]*

Rules: - Must start with letter or underscore - Can contain letters, digits, underscores
- Case sensitive - Cannot be keywords - Maximum length: 255 characters

3. Literals

Integer Literals

DECIMAL_INT := [0-9] ('_'? [0-9])*
HEX_INT := 0[xX] [0-9a-fA-F] ('_'? [0-9a-fA-F])*
BINARY_INT := 0[bB] [01] ('_'? [01])*
OCTAL_INT := 0[oO] [0-7] ('_'? [0-7])*

INT_SUFFIX := (i8|i16|i32|i64|u8|u16|u32|u64)

INTEGER := (DECIMAL_INT | HEX_INT | BINARY_INT | OCTAL_INT)
INT_SUFFIX?

Rules: - Underscores are digit separators and may appear only *between* two digits — never leading, trailing, or consecutive (1__2 and 12_ are invalid). - A suffix fixes the literal's type: i8→Int8, i16→Int16, i32→Int32, i64→Int64, u8→UInt8, u16→UInt16, u32→UInt32, u64→UInt64. A suffixed literal whose value does not fit the named type is a compile-time error.

Examples:

42
1_000_000
0xFF_FF
0b1010_1010
0o755
42i32
255u8

Floating Point Literals

FLOAT_BODY := DECIMAL_INT '.' [0-9] ('_'? [0-9])* EXPONENT?
| DECIMAL_INT EXPONENT
EXPONENT := [eE] [+-]? [0-9] ('_'? [0-9])*

FLOAT_SUFFIX := (f32|f64)

`FLOAT := FLOAT_BODY FLOAT_SUFFIX?`

Rules: - The underscore rules for integer literals apply (separators between digits only). - A suffix fixes the literal's type: `f32`→`Float32`, `f64`→`Float64`.

Examples:

```
3.14
1.0e10
2.5e-3
42.0f32
```

String Literals

`STRING := ''' (CHAR | ESCAPE_SEQUENCE)* '''`
`RAW_STRING := 'r''' .*? '''`

`ESCAPE_SEQUENCE := '\' (n|t|r|0|\\|'|"|x[0-9a-fA-F]{2}|u{[0-9a-fA-F]{1,6}})`

Examples:

```
"Hello, World!"
"Line 1\nLine 2"
"\x41\x42\x43"
"\u{1F600}"
r"Raw string with \n literal backslashes"
```

Character Literals

`CHAR := '\'' (CHAR_CONTENT | ESCAPE_SEQUENCE) '\''`
`CHAR_CONTENT := [^'\\]`

Examples:

```
'a'
'\n'
'\x41'
'\u{1F600}'
```

Boolean Literals

`BOOL := true | false`

4. Operators

Arithmetic

`+ - * / % **`

Comparison

`== != < <= > >=`

Logical

`&& || !`

Bitwise

`& | ^ ~ << >>`

Assignment

`= += -= *= /= %= **= &= |= ^= <<= >>=`

Range

`.. ..=`

Other

`? :: . -> =>`

Notes: - `?` is the try operator (postfix). There is no ternary conditional operator; `if` is an expression. - `:` appears only as a delimiter (type annotations, struct fields), not as an operator.

5. Delimiters

`( ) [ ] { } , ; :`

6. Comments

```
// Single line comment
/* Multi-line comment */
/** Documentation comment */
```

Rules: - Single line comments start with `//` and continue to end of line - Multi-line comments start with `/*` and end with `*/` - Multi-line comments can be nested - Documentation comments start with `/**` and are associated with following declaration

7. Whitespace

- Space (U+0020)
- Tab (U+0009)
- Newline (U+000A)
- Carriage Return (U+000D)

Whitespace separates tokens and is otherwise ignored. Statement termination is defined by semicolons in the syntax grammar.

8. Token Precedence

When multiple token patterns could match: 1. Keywords take precedence over identifiers 2. Longer operators take precedence over shorter ones 3. Comments are ignored in token stream

9. Reserved Tokens

Reserved for future use (recognized as keywords but not used by any Core v1 grammar production):

async, await, yield, where, macro, unsafe, extern, import, export,
null,
and, or, not, is, dyn

Lexical Analysis Rules

1. **Maximal Munch:** Always match the longest possible token
2. **Whitespace:** Ignored except for token separation
3. **Encoding:** Source files must be valid UTF-8

Error Handling

Invalid tokens should produce specific error messages: - Invalid character: “Unexpected character ‘X’ at line Y:Z” - Unterminated string: “Unterminated string literal at line Y:Z” - Invalid number: “Invalid number format at line Y:Z” ##
Conformance A conforming Core v1 implementation **MUST** follow the requirements in this document. Any deviations or extensions **MUST** be explicitly documented by the implementation.

STARK Syntax Grammar Specification

Overview

This document defines the concrete syntax grammar for STARK using Extended Backus-Naur Form (EBNF).

Grammar Notation

`::=` means "is defined as"
`|` means "or" (alternative)
`?` means "optional" (zero or one)
`*` means "zero or more"
`+` means "one or more"
`()` means "grouping"
`[]` means "optional"
`{}` means "zero or more repetitions"

Top-Level Grammar

Program

`Program ::= Item*`

`Item ::= Visibility? (Function
| Struct
| Enum
| Trait
| Impl
| Const
| TypeAlias
| Use
| Module)`

`Visibility ::= 'pub' | 'priv'`

Generic Parameters and Arguments

`GenericParams ::= '<' GenericParam (',' GenericParam)* ','? '>'`

`GenericParam ::= IDENTIFIER (':' TraitBounds)?`

`TraitBounds ::= TraitBound ('+' TraitBound)*`

`TraitBound ::= Path GenericArgs?`

`GenericArgs ::= '<' GenericArg (',' GenericArg)* ','? '>'`

`GenericArg ::= Type
| IDENTIFIER '=' Type // Associated type
binding, e.g. Iterator<Item = T>`

Notes: - Generic parameters may appear on functions, structs, enums, traits, impl blocks, and type aliases. - Generic arguments in *expressions* are inferred at the use site. The only explicit form is `path::<Args>` on a path expression (see `PathExpression`), for calls where inference is impossible, e.g. `size_of::()`.

Function Definition

Function ::= FunctionSig Block

FunctionSig ::= 'fn' IDENTIFIER GenericParams? '(' ParameterList? ')' ('->' ReturnType)?

ReturnType ::= Type | '!' // '!' is the never type (diverging function)

ParameterList ::= Receiver (',' Parameter)* ','? | Parameter (',' Parameter)* ','?

Receiver ::= 'self' | '&' 'self' | '&' 'mut' 'self'

Parameter ::= IDENTIFIER ':' Type | 'mut' IDENTIFIER ':' Type

Block ::= '{' Statement* Expression? '}'

Block semantics: - The optional final Expression has no terminating semicolon and is the block's value; a block without one evaluates to Unit. - Disambiguation: at each position inside a block, the parser first attempts to parse a Statement; a trailing Expression is recognized only immediately before }. An expression followed by ; is always a statement.

Notes: - A receiver is only valid on functions declared inside a trait or impl block. - A function without a -> annotation returns Unit. Return types are never inferred.

Data Type Definitions

Struct ::= 'struct' IDENTIFIER GenericParams? '{' FieldList? '}'

FieldList ::= Field (',' Field)* ','?

Field ::= IDENTIFIER ':' Type | 'pub' IDENTIFIER ':' Type

Enum ::= 'enum' IDENTIFIER GenericParams? '{' VariantList? '}'

VariantList ::= Variant (',' Variant)* ','?

Variant ::= IDENTIFIER | IDENTIFIER '(' TypeList ')' | IDENTIFIER '{' FieldList '}'

Trait ::= 'trait' IDENTIFIER GenericParams? '{' TraitItem* '}'

TraitItem ::= FunctionSig ';' // Required method (signature only) | FunctionSig Block // Method with default

```

body
    | AssociatedType

AssociatedType ::= 'type' IDENTIFIER ';'

Impl ::= 'impl' GenericParams? Type '{' ImplItem* '}'
      | 'impl' GenericParams? TraitBound 'for' Type '{' ImplItem*
      '}'

ImplItem ::= Visibility? Function
          | 'type' IDENTIFIER '=' Type ';' // Associated type
value

```

Module Definition

```

Module ::= 'mod' IDENTIFIER ';'
        | 'mod' IDENTIFIER ModuleBlock

ModuleBlock ::= '{' Item* '}'

```

Type Alias

```

TypeAlias ::= 'type' IDENTIFIER GenericParams? '=' Type ';'

```

Statements

```

Statement ::= ';' // Empty statement
          | Expression ';' // Expression statement
          | BlockExpression ';' // Block-formed expression
statement
    | 'let' LetStatement
    | 'return' Expression? ';'
    | 'break' Expression? ';'
    | 'continue' ';'

BlockExpression ::= Block
                | IfExpression
                | MatchExpression
                | LoopExpression

LetStatement ::= IDENTIFIER ':' Type ';'
              | 'mut' IDENTIFIER ':' Type ';'
              | IDENTIFIER ':' Type '=' Expression ';'
              | IDENTIFIER '=' Expression ';'
              | 'mut' IDENTIFIER ':' Type '=' Expression ';'
              | 'mut' IDENTIFIER '=' Expression ';'

```

Block-formed expression statements: - An expression whose outermost form is a block, if, match, loop, while, or for may stand as a statement without a terminating semicolon (if `ready { start(); }` mid-block is a statement). - Statement parsing is greedy: when a block-formed expression begins at statement position and is not followed by `;`, it is complete as a statement — a following token

does NOT extend it into a larger expression (`{ if c { } - 1; }` is an if statement followed by an error, not a subtraction). Parenthesize to use the value: `(if c { 1 } else { 2 }) - 1`. - Exception (trailing expression): a block-formed expression immediately before the closing `}` of its block and not followed by `;` is the block's trailing Expression (its value), not a statement — so `fn max(...) -> T { if a > b { a } else { b } }` returns the if's value.

Expressions

Expression ::= AssignmentExpression

AssignmentExpression ::= RangeExpression
| RangeExpression AssignOp

AssignmentExpression

AssignOp ::= '=' | '+=' | '-=' | '*=' | '/=' | '%=' | '**='
| '&=' | '|=' | '^=' | '<=>' | '>=>'

Place Expressions

The left-hand side of an assignment must be a *place expression* — an expression that denotes a memory location. Place expressions are:

PlaceExpression ::= Path // Variable
| PlaceExpression '.' IDENTIFIER // Field
| PlaceExpression '.' INTEGER // Tuple
field
| PlaceExpression '[' Expression ']' // Index
| '*' Expression //
Dereference
| '(' PlaceExpression ')'

Assignment to any other expression form is a compile-time error (checked semantically; the grammar accepts any expression on the left).

RangeExpression ::= LogicalOrExpression (('..' | '..=')
LogicalOrExpression)?

LogicalOrExpression ::= LogicalAndExpression ('||'
LogicalAndExpression)*

LogicalAndExpression ::= EqualityExpression ('&'
EqualityExpression)*

EqualityExpression ::= RelationalExpression (('==' | '!=')
RelationalExpression)?

RelationalExpression ::= BitwiseOrExpression ('<' | '<=' | '>' |
'>=') BitwiseOrExpression)?

BitwiseOrExpression ::= BitwiseXorExpression ('|'
BitwiseXorExpression)*

```

BitwiseXorExpression ::= BitwiseAndExpression ('^'
BitwiseAndExpression)*

BitwiseAndExpression ::= ShiftExpression ('&' ShiftExpression)*

ShiftExpression ::= AdditiveExpression (('<<' | '>>')
AdditiveExpression)*

AdditiveExpression ::= MultiplicativeExpression (('+' | '-')
MultiplicativeExpression)*

MultiplicativeExpression ::= ExponentiationExpression (('*' | '/'
| '%') ExponentiationExpression)*

ExponentiationExpression ::= CastExpression ('**'
ExponentiationExpression)?

CastExpression ::= UnaryExpression ('as' Type)*

UnaryExpression ::= PostfixExpression
                    | '-' UnaryExpression
                    | '!' UnaryExpression
                    | '~' UnaryExpression
                    | '&' 'mut'? UnaryExpression // Reference
(immutable or mutable borrow)
                    | '*' UnaryExpression        // Dereference

PostfixExpression ::= PrimaryExpression
                    | PostfixExpression '(' ArgumentList?
')'        // Function call
                    | PostfixExpression '[' Expression
']'        // Index (or slice, when the index is a range)
                    | PostfixExpression '.'
IDENTIFIER // Field access / method reference
                    | PostfixExpression '.'
INTEGER    // Tuple access
                    | PostfixExpression
'?'        // Try operator

ArgumentList ::= Expression (',' Expression)* ',' '?'

PrimaryExpression ::= PathExpression
                    | Literal
                    | '(' Expression ')' // Grouping
                    | TupleLiteral
                    | ArrayLiteral
                    | StructLiteral
                    | IfExpression
                    | MatchExpression
                    | LoopExpression
                    | Block

PathExpression ::= Path (':' GenericArgs)? // e.g. x,
String::from, Color::Red, size_of::<Int32>

```

Notes: - `path::<Args>` (explicit generic arguments on a path expression) is permitted only when type inference cannot determine the arguments, e.g. `size_of::<Int32>()`. It is the only form of expression-level generic arguments in Core v1. - Chained comparisons (`a < b < c`, `a == b == c`) are syntax errors: the relational and equality productions are non-associative (at most one operator). Parenthesize to compare a `Bool` result explicitly.

Control Flow Expressions

`IfExpression ::= 'if' Expression Block ('else' 'if' Expression Block)* ('else' Block)?`

`MatchExpression ::= 'match' Expression '{' MatchArm* '}'`

`MatchArm ::= Pattern '=>' Expression ',' '?'`

`Pattern ::= Literal`
 | `'_'` // Wildcard
 | `IDENTIFIER` // Binding (or
unit variant/const, see note)
 | `Path` // Unit enum
variant or const, e.g. `Color::Red`
 | `Path '(' PatternList? ')'` // Tuple enum
variant, e.g. `Option::Some(x)`
 | `Path '{' FieldPatternList? '}'` // Struct or
struct enum variant
 | `'(' PatternList? ')'` // Tuple
 | `'[' PatternList? ']'` // Array

`PatternList ::= Pattern (',' Pattern)* ',' '?'`

`FieldPatternList ::= FieldPattern (',' FieldPattern)* ',' '?'`

`FieldPattern ::= IDENTIFIER ':' Pattern`
 | `IDENTIFIER` // Shorthand:
binds field to same-named variable

`LoopExpression ::= 'loop' Block`
 | `'while' Expression Block`
 | `'for' IDENTIFIER 'in' Expression Block`

Note on pattern name resolution: a single `IDENTIFIER` pattern that resolves to a unit enum variant or a constant in scope matches by value; otherwise it introduces a new binding. Multi-segment `Path` patterns always match by value.

Literals

`Literal ::= INTEGER`
 | `FLOAT`
 | `STRING`
 | `CHAR`
 | `BOOLEAN`

```

TupleLiteral ::= '(' ' ' )' // Unit
value
                | '(' Expression ',' ' ' )' //
Single-element tuple
                | '(' Expression (',' Expression)+ ',' '?' ' ' )' // N-
element tuple

ArrayLiteral ::= '[' ExpressionList? ']'
                | '[' Expression ';' Expression ']' //
Repeat: [value; count]

ExpressionList ::= Expression (',' Expression)* ',' '?'

StructLiteral ::= Path '{' FieldInitList? '}'

FieldInitList ::= FieldInit (',' FieldInit)* ',' '?'

FieldInit ::= IDENTIFIER ':' Expression
              | IDENTIFIER // Shorthand: field:
field

```

Notes: - (expr) is grouping; (expr,) is a single-element tuple; () is the unit value.
The trailing comma distinguishes the 1-tuple from grouping. - In [value; count],
count must be a compile-time constant expression of an unsigned integer type, and
value's type must implement Copy (the value is copied count times).

Struct Literal Restriction

To avoid ambiguity with block-taking constructs, a struct literal may NOT appear as
the outermost expression in: - the condition of an if or while, - the scrutinee of a
match, - the iterable of a for.

Wrap the struct literal in parentheses in these positions:

```
if (Config { verbose: true }).verbose { do_thing(); }
```

Types

```

Type ::= PrimitiveType
        | Path GenericArgs? // Named type, possibly
generic: Vec<Int32>, Option<T>
        | '[' Type ';' INTEGER ']' // Array type
        | '[' Type ']' // Slice type (unsized;
used behind references)
        | TupleType
        | '(' Type ')' // Grouping (the type
T, not a tuple)
        | '&' Type // Reference type
        | '&' 'mut' Type // Mutable reference
type
        | 'fn' '(' TypeList? ')' ('->' Type)? // Function type
(non-capturing)

```

```

TupleType ::= '(' ')' // Unit type
           | '(' Type ',' ')' // Single-element tuple
type
           | '(' Type (',' Type)+ ',' '?' ')' // N-element tuple
type

```

```

PrimitiveType ::= 'Int8' | 'Int16' | 'Int32' | 'Int64'
               | 'UInt8' | 'UInt16' | 'UInt32' | 'UInt64'
               | 'Float32' | 'Float64'
               | 'Bool' | 'Char' | 'String' | 'Unit' | 'str'

```

```

TypeList ::= Type (',' Type)* ',' '?'

```

Notes: - A function type without $\rightarrow$ returns Unit. - The never type ! may appear only as a function return type (see Return Type). - (T) is the type T (grouping); the single-element tuple type is (T,), mirroring TupleLiteral on the value side.

Other Constructs

```

Const ::= 'const' IDENTIFIER ':' Type '=' Expression ';'

```

```

Use ::= 'use' UseTree ';'

```

```

UseTree ::= Path ('as' IDENTIFIER)?
          | Path '::' '*'
          | Path '::' 'self'
          | Path '::' '{' UseTreeList? '}'

```

```

UseTreeList ::= UseTree (',' UseTree)* ',' '?'

```

```

Path ::= PathSegment ('::' PathSegment)*
PathSegment ::= IDENTIFIER | PrimitiveType | 'self' | 'Self' |
'super' | 'crate'

```

Notes: - Self is valid only inside a trait or impl block, and only as the *first* segment of a path. As a complete one-segment path in type position it denotes the implementing type; with further segments it projects an associated item, e.g. Self::Item. - self, super, and crate are likewise valid only as leading segments (self/super may repeat at the front: super::super::x). - A PrimitiveType keyword is valid only as the *first* segment, where it names the primitive's associated items: String::from(s). (Primitive type names are keywords in the lexical grammar, so without this rule such calls would be unparseable.)

Operator Precedence (Highest to Lowest)

1. Primary expressions, field access, array access, function calls, try operator (?)
2. Unary operators (-, !, ~, &, &mut, *)
3. Cast (as)
4. Exponentiation (**)
5. Multiplicative (*, /, %)

6. Additive (+, -)
7. Shift (<<, >>)
8. Bitwise AND (&)
9. Bitwise XOR (^)
10. Bitwise OR (|)
11. Relational (<, <=, >, >=)
12. Equality (==, !=)
13. Logical AND (&&)
14. Logical OR (||)
15. Range (.., ..=)
16. Assignment (=, +=, -=, etc.)

Associativity

- Left associative: Most binary operators
- Right associative: Assignment operators, exponentiation
- Non-associative: Comparison operators, range operators

Statement vs Expression

- Statements do not return values and end with semicolons (block-formed expression statements may omit the semicolon — see Statements)
- Expressions return values and can be used as statements with semicolons
- Blocks are expressions (return value of last expression, or Unit if none)
- Control flow constructs are expressions

Whitespace and Semicolons

- Semicolons are required to terminate statements, except after block-formed expression statements (see Statements)
- Semicolons are optional after the last expression in a block
- Newlines are not significant except for line comments
- Trailing commas are allowed in lists

Parsing Notes

- When a > is expected in generic-argument position and the next token is >>, >=, or = (maximal munch), the parser MUST split off a single > and re-tokenize the remainder (>, >=, or = respectively) — so `Vec<Vec<Int32>>`, `Vec<Vec<Int32>>= v`, and `Vec<Int32>= v` all parse.
- Nested tuple-field access `pair.0.1` lexes the `0.1` as a FLOAT token (maximal munch); after `.`, the parser MUST split an unsuffixed `digits-.-digits` FLOAT into two INTEGER tuple indices.
- In expression position, a bare < after an identifier is always the relational operator; explicit generic arguments require the `::<` form (turbofish), so no lookahead disambiguation is required.

Informative Future Grammar Sketches (Non-Normative)

Reserved grammar constructs for later implementation:

```
// Async functions (future)
AsyncFunction ::= 'async' 'fn' IDENTIFIER '(' ParameterList? ')'
('->' Type)? Block

// Lambda expressions / capturing closures (future)
Lambda ::= '|' ParameterList? '|' (Expression | Block)

// Open-ended ranges (future)
OpenRange ::= Expression '..' | '..' Expression | '..'

// Lifetime annotations (future)
LifetimeParam ::= '\' IDENTIFIER
```

Conformance

A conforming Core v1 implementation **MUST** follow the requirements in this document. Any deviations or extensions **MUST** be explicitly documented by the implementation.

STARK Type System Specification

Overview

STARK features a static type system with type inference, ownership semantics, and memory safety guarantees.

Core Principles

1. **Static Typing:** All types resolved at compile time
2. **Type Inference:** Types can be inferred when unambiguous
3. **Memory Safety:** No null pointer dereferences, no use-after-free
4. **Ownership:** Clear ownership and borrowing rules
5. **Zero-cost Abstractions:** Type safety without runtime overhead

Primitive Types

Integer Types

```
Int8    // 8-bit signed integer (-128 to 127)
Int16   // 16-bit signed integer (-32,768 to 32,767)
Int32   // 32-bit signed integer (-2^31 to 2^31-1)
Int64   // 64-bit signed integer (-2^63 to 2^63-1)

UInt8   // 8-bit unsigned integer (0 to 255)
UInt16  // 16-bit unsigned integer (0 to 65,535)
UInt32  // 32-bit unsigned integer (0 to 2^32-1)
UInt64  // 64-bit unsigned integer (0 to 2^64-1)
```

Default integer type is Int32 for literals that fit, Int64 otherwise.

Floating Point Types

```
Float32 // 32-bit IEEE 754 floating point
Float64 // 64-bit IEEE 754 floating point
```

Default floating point type is Float64.

Other Primitive Types

```
Bool    // Boolean: true or false
Char    // Unicode scalar value (32-bit)
Unit    // Unit type: () – represents no meaningful value
str     // String slice (unsized), typically used behind
references: &str
```

String Type

```
String // UTF-8 encoded string, heap-allocated, growable
```

String Slice Type

```
// str is an unsized string slice type. It is used via references.
let s: &str = "hello";
let owned: String = String::from(s);
```

Never Type

! (the never type) is the type of expressions that never produce a value, such as calls to panic. It may appear only as a function return type.

```
fn panic(message: &str) -> !
```

Rules: - An expression of type ! coerces to any other type. This allows, for example, `panic(...)` to appear in one arm of an `if` or `match` whose other arms produce a value. - A function returning ! MUST NOT return normally.

Composite Types

Array Types

```
[T; N]    // Fixed-size array of N elements of type T (sized)
[T]       // Slice of elements of type T (unsized)
```

A slice `[T]` is an *unsized* view into a contiguous sequence (an array, or the elements of a `Vec<T>`). Like `str`, it is used behind references:

```
let fixed: [Int32; 5] = [1, 2, 3, 4, 5];
let view: &[Int32] = &fixed;           // Array reference coerces
to slice reference
let part: &[Int32] = &fixed[1..4];     // Slicing with a range
yields &[T]
```

A local variable cannot have the bare unsized type `[T]`; use `[T; N]` for a fixed-size array or `Vec<T>` (standard library) for a growable one.

Indexing rules: - `expr[i]` where `i` is an integer denotes a *place* of type `T` (bounds-checked; traps on out-of-bounds). - `expr[r]` where `r` is a `Range` denotes a *place* of the unsized slice type `[T]`; traps if the range is out of bounds or inverted. Because `[T]` is unsized, a slice place must be borrowed immediately: `&expr[r]` has type `&[T]` and `&mut expr[r]` has type `&mut [T]`.

Range Type

The range operators produce values of the standard library `Range<T>` type:

```
let r = 0..10;      // Range<Int32>: 0,1,...,9 (half-open)
let ri = 0..=9;     // RangeInclusive<Int32>: 0,1,...,9
```

Ranges over integer types implement `Iterator` and may be used with `for` loops and slicing.

Tuple Types

```
()          // Empty tuple (Unit type)
(T,)        // Single-element tuple
(T1, T2)    // Two-element tuple
(T1, T2, T3) // Three-element tuple
// ... up to reasonable limit (e.g., 16 elements)
```

Examples:

```
let empty: () = ();
let single: (Int32,) = (42,);
let pair: (Int32, String) = (42, "hello");
```

Struct Types

```
struct Point {
    x: Float64,
    y: Float64
}

struct Person {
    name: String,
    age: Int32,
    pub email: String // Public field
}
```

Restriction (Core v1): struct and enum declarations **MUST NOT** declare fields whose *written* type is a reference type (`&T`, `&mut T`). Instantiating a generic type with a reference type argument (e.g. `Option<&T>`) is permitted; the resulting value is *borrow-carrying* — see “References and Lifetimes” below.

Enum Types

```
enum Color {
    Red,
    Green,
    Blue
}

enum Option<T> {
    Some(T),
    None
}

enum Result<T, E> {
    Ok(T),
    Err(E)
}
```

Reference Types

```
&T        // Immutable reference to T
&mut T    // Mutable reference to T
```

Function Types

Function types are written with the `fn` keyword and denote *non-capturing* functions (named functions or function references). Capturing closures are a future extension.

```
fn(T1, T2) -> R    // Function taking T1, T2 and returning R
fn() -> R          // Function taking no parameters, returning R
fn(T)              // Function taking T, returning Unit
```

Type Aliases

```
type Age = Int32;
type Point2D = (Float64, Float64);
type ErrorCode = Int32;
```

The Self Type

Within a trait or impl block, Self refers to the implementing type:

```
trait Eq {
    fn eq(&self, other: &Self) -> Bool;
}
```

Ownership and Borrowing

Ownership Rules

1. Each value has exactly one owner
2. When the owner goes out of scope, the value is dropped
3. Values can be moved (ownership transfer) or borrowed (temporary access)

Move Semantics

```
let a = String::from("hello");
let b = a; // a is moved to b, a is no longer valid
// print(a); // Error: use of moved value
```

Borrowing Rules

1. You can have either one mutable reference or any number of immutable references
2. References must always be valid (no dangling pointers)
3. References cannot outlive the data they refer to

```
fn borrow_immutable(s: &String) {
    // Can read but not modify
}
```

```
fn borrow_mutable(s: &mut String) {
    // Can read and modify
}
```

```
let mut text = String::from("hello");
```

```
borrow_immutable(&text);      // OK
borrow_mutable(&mut text);    // OK
```

References and Lifetimes (Core v1)

Core v1 has no lifetime annotations. Instead, it enforces the following conservative rules, which are normative:

1. **No declared reference fields.** Struct, enum variant, and tuple struct declarations MUST NOT write a reference type (`&T`, `&mut T`) as a field type. (Generic *instantiation* with a reference argument is allowed; see Borrow-Carrying Types below.)
2. **References derive from parameters.** A function may return a reference (or borrow-carrying value) only if every control-flow path derives it from one of its reference parameters (directly, by field/index projection, or by calling a function that itself obeys this rule). Returning a reference derived from a local variable or a by-value parameter is a compile-time error.
3. **Shortest-input-lifetime rule.** The lifetime of a returned reference (or borrow-carrying value) is the *shortest* of the lifetimes of all reference parameters from which it could have been derived. Callers MUST treat it as invalidated as soon as the shortest-lived of those arguments is invalidated.
4. **Local borrows are lexically scoped.** A borrow bound to a variable (`let r = &x;`) lasts from its creation to the end of its enclosing block (or until the borrower is explicitly dropped with `drop(...)`). The borrow checker enforces the exclusive/shared rules over that entire region — even if the reference is never used again. To end such a borrow early, introduce an inner block or call `drop`. A *temporary* borrow that is not bound to a variable (e.g. `f(&x);`) ends at the end of its enclosing statement, so `f(&x); g(&mut x);` is legal.

```
let mut s = String::from("hello");
{
    let r = &s;          // Borrow begins
    println(r.as_str());
}                        // Borrow ends with the block
s.push('!');            // OK: no live borrow
```

Example of rule 3:

```
fn longest(x: &str, y: &str) -> &str {
    if x.len() > y.len() { x } else { y }
}
// The returned reference is valid only while BOTH x's and y's
// referents
// are valid, regardless of which branch was taken.
```

These rules reject some safe programs (that is intentional). Lifetime annotations that relax rule 3, user structs with declared reference fields, and non-lexical (last-use) borrow regions are future extensions; because the lexical rule is strictly more conservative, adopting last-use regions later cannot break conforming programs.

Borrow-Carrying Types

A type is **borrow-carrying** if it is: - a reference type `&T` or `&mut T`; or - a generic type with at least one borrow-carrying type argument (e.g. `Option<&T>`, `(Int32, &str)`); or - an implementation-provided borrowed view type (the standard library's borrowed iterators such as `VecIter<T>`, `CharsIter`, `KeysIter<K>`, and the slice types behind references).

A borrow-carrying value is treated *as if it were a reference* for all rules in this section: - It counts as a live borrow of the value(s) it was derived from — shared for `&`-derived values, exclusive for `&mut`-derived values — for the borrow region defined by rule 4. - It MUST NOT be stored in a user-declared struct or enum field, returned except under rules 2–3, or otherwise escape the lifetime of its source. - Binding it with `let` is permitted; the binding is subject to rule 4.

This is a taint rule: borrow-carrying-ness propagates through generic instantiation and function returns, and the checker tracks the source variables each borrow-carrying value may borrow from. It is what makes APIs like `Vec::get(&self, i) -> Option<&T>` sound without lifetime annotations: the returned `Option<&T>` is an active shared borrow of the `Vec` until it goes out of scope.

Type Inference

Local Type Inference

Types of `let` bindings are inferred from their initializers. Inference is local: it uses the initializer expression and, where necessary, later uses within the same function body. Function signatures are always fully explicit, so inference never crosses function boundaries.

```
let x = 42;           // Inferred as Int32
let y = 3.14;         // Inferred as Float64
let z = [1, 2, 3];    // Inferred as [Int32; 3]
```

Function Return Types

Function return types are NOT inferred. A function without a `->` annotation returns `Unit`.

```
fn add(a: Int32, b: Int32) -> Int32 {
    a + b
}

fn greet(name: &str) {           // Returns Unit
    println(name);
}
```


Generic Type Inference

Generic arguments at call sites are inferred from argument and expected types:

```
fn identity<T>(x: T) -> T {  
    x  
}  
  
let result = identity(42); // T inferred as Int32
```

Subtyping and Coercion

Numeric Coercions

No implicit numeric conversions. Explicit casting required:

```
let x: Int32 = 42;  
let y: Int64 = x as Int64; // Explicit cast required
```

Cast semantics (as): - Between integer types: value-preserving if in range; a cast whose value does not fit the target type is a runtime error and MUST trap. - Between floating point types: rounds to nearest representable value. - Integer to float and float to integer: float-to-int truncates toward zero and MUST trap if the result does not fit the target type (including NaN/Inf).

Reference Coercions

```
&mut T -> &T           // Mutable reference to immutable reference  
&T -> &T               // Same type (identity)
```

Array to Slice Coercion

```
&[T; N] -> &[T]        // Array reference to slice reference  
&mut [T; N] -> &mut [T] // Mutable array reference to mutable  
slice reference
```

Never Coercion

```
! -> T                // The never type coerces to any type
```

Trait System (Basic)

Trait Definition

```
trait Display {  
    fn fmt(&self) -> String;  
}
```

```
trait Eq {
    fn eq(&self, other: &Self) -> Bool;
}
```

Trait Implementation

```
impl Display for Int32 {
    fn fmt(&self) -> String {
        // Convert integer to string
    }
}

impl Eq for Point {
    fn eq(&self, other: &Point) -> Bool {
        self.x == other.x && self.y == other.y
    }
}
```

Associated Types

Traits may declare associated types; implementations assign them:

```
trait Iterator {
    type Item;
    fn next(&mut self) -> Option<Self::Item>;
}

struct Counter { count: Int32 }

impl Iterator for Counter {
    type Item = Int32;
    fn next(&mut self) -> Option<Int32> {
        self.count += 1;
        Some(self.count)
    }
}
```

Bounds may constrain associated types with bindings: `I: Iterator<Item = T>`.

Type Checking Rules

Assignment Compatibility

`let x: T = expr; // expr must have type T or be coercible to T`

Function Call Compatibility

```
fn f(param: T) { ... }
f(arg); // arg must have type T or be coercible to T
```

Arithmetic Operations

```
// Binary arithmetic requires same types
let result = x + y; // x and y must have the same numeric type

// Comparison operators
let cmp = x < y; // x and y must have the same comparable type
```

Logical Operations

```
let both = a && b; // a and b must be Bool
let negated = !x; // x must be Bool
```

For Loops

for x in expr requires expr to have a type that implements Iterator; the loop variable binds successive Item values. Integer ranges implement Iterator; slices and collections provide .iter() methods (see the standard library specification).

```
for i in 0..10 {
    println(i.fmt());
}
```

Control Flow Typing

- **if/else**: an if with an else is an expression whose branches must unify to a single type (identical types, or unifiable via the ! coercion — e.g. one branch may panic). An if *without* else has type Unit, and its branch must also have type Unit.
- **match**: all arm expressions must unify to a single type (with ! coercion permitted per arm).
- **loop**: a loop expression's type is the unified type of the operands of its break value; statements; if every break is bare (no value) or the loop has no break, the type is Unit (a loop with no break has type !).
- **while / for**: always have type Unit. break inside while/for MUST NOT carry a value.
- **Block**: the type of the trailing expression, or Unit if there is none.

Method Calls and Auto-Borrowing

A method call recv.m(args) is resolved as follows:

1. **Candidate collection.** Candidates are, in priority order:
 1. inherent methods — fn m in impl RecvType { ... } blocks;
 2. trait methods — fn m from any trait that RecvType implements, where the trait is *in scope* (defined in, or imported via use into, the current module, or in the prelude). Inherent methods shadow trait methods of the same name.
2. **Ambiguity.** If two or more traits in scope supply an applicable m and no inherent method exists, the call is a compile-time error. Disambiguate with a fully-qualified call, passing the receiver explicitly: TraitName::m(&recv, args).

3. **Receiver coercion (auto-borrowing).** Let S be the receiver expression's type. Candidates are matched trying these receiver forms in order: `self: S` (by value), `self: &S` (auto-borrow: the compiler inserts `&`), `self: &mut S` (auto-mutable-borrow: the compiler inserts `&mut`; requires the receiver to be a mutable place). Auto-borrows follow the normal borrow rules.
4. **Auto-dereference.** If S is `&T` or `&mut T` and no candidate matches on S , resolution retries with T (one level of dereference, applied repeatedly for nested references). Combined with step 3 this allows `(&v).len()` and `v.len()` to resolve identically.
5. **Visibility.** Private methods are callable only per the module rules (07-Modules-and-Packages.md).

Trait methods can always be called in fully-qualified function form — `Display::fmt(&x)` — since a method is an ordinary function whose first parameter is the receiver.

Operators and Traits

Operator expressions on **primitive types** have built-in meaning (Numeric Semantics below). On **generic type parameters**, operators desugar to trait method calls, so the corresponding bound is required:

Operator	Requires bound	Desugars to
<code>==, !=</code>	<code>T: Eq</code>	<code>Eq::eq(&a, &b)</code> (negated for <code>!=</code>)
<code><, <=, >, >=</code>	<code>T: Ord</code>	<code>Ord::cmp(&a, &b)</code> compared to <code>Ordering</code>
<code>+ - * / % **</code> , bitwise, shifts	<code>T: Num</code>	primitive operation after monomorphization

`Num` is a compiler-known marker trait implemented by exactly the built-in numeric types (`Int8–Int64`, `UInt8–UInt64`, `Float32`, `Float64`); user types cannot implement it in Core v1. Operator overloading for user-defined types (`Add/Sub/...` traits) is a future extension. `&&`, `||`, and `!` (on `Bool`) are built-in, short-circuiting, and not overloadable.

```
fn max<T: Ord>(a: T, b: T) -> T {
    if a > b { a } else { b }    // a > b desugars to Ord::cmp(&a,
                                // &b)
}
```

Evaluation Order (Core v1)

Runtime evaluation order is defined solely by `CORE-V1-ABSTRACT-MACHINE.md` (`EXEC-EVAL-001`, `EXEC-ONCE-001`, and `EXEC-ASSIGN-001`). Type checking must preserve those rules and must not introduce an evaluation or ownership transfer merely to select a type, method, or trait implementation.

Copy and Drop (Soundness Rules)

- Copy may be implemented for a type only if **all** of its fields are Copy. Violations are compile-time errors.
- A type **MUST NOT** implement both Copy and Drop (a copyable type would run its destructor once per copy).
- `Drop::drop` **MUST NOT** be called explicitly; use the free function `drop(value)`. After `drop(v)`, `v` is moved-from.
- **Partial moves**: moving a field out of a struct is permitted only if the struct's type does not implement Drop. After a partial move, the whole value may no longer be used or moved, but its remaining fields may be.
- **Reinitialization**: assigning a new value to a moved-from `let mut` variable makes it valid again (definite-assignment tracking).
- Moving an element out of an indexed place (`v[i]`) is a compile-time error; use APIs that transfer ownership explicitly (e.g. `Vec::remove`, `Vec::pop`).

The execution and ordering of moves, copies, replacement, `drop(value)`, partial-value destruction, and trap termination are defined solely by `CORE-V1-ABSTRACT-MACHINE.md`.

Error Types and Handling

Option Type

```
enum Option<T> {  
    Some(T),  
    None  
}
```

Result Type

```
enum Result<T, E> {  
    Ok(T),  
    Err(E)  
}
```

Error Propagation

```
fn might_fail() -> Result<Int32, String> {  
    // ...  
}  
  
fn caller() -> Result<Int32, String> {  
    let value = might_fail()?; // Early return on error  
    Ok(value * 2)  
}
```

Try Operator (?) Typing

The try operator is defined for `Result<T, E>` and `Option<T>`: - If `expr` has type `Result<T, E>`, then `expr?` has type `T` and propagates `Err(E)` to the nearest enclosing function returning `Result<_, E>`. - If `expr` has type `Option<T>`, then `expr?` has type `T` and propagates `None` to the nearest enclosing function returning `Option<_>`.

The enclosing function's return type must be compatible with the propagated type.

Generics (Core v1)

Core v1 supports parametric polymorphism for functions, structs, enums, and traits.

Rules: - Generic parameters are introduced with `<T, U, ...>` after the item name. - Generic parameters are in scope within the item body and signatures. - All generic parameters used in an item MUST be declared by that item. - Instantiation occurs at use sites; Core v1 permits monomorphization or dictionary-passing, but the observable behavior MUST be equivalent. - Generic arguments in expressions are inferred. When inference is impossible (no parameter or usable result type mentions the type parameter), explicit arguments are supplied with the turbofish form on a path expression: `size_of::<Int32>()`. This is the only expression-level generic-argument syntax in Core v1.

```
struct Pair<T> {
    first: T,
    second: T
}

fn max<T: Ord>(a: T, b: T) -> T {
    if a > b { a } else { b }
}
```

Trait Bounds

```
fn max<T: Ord>(a: T, b: T) -> T { if a > b { a } else { b } }
```

Rules: - A bound `T: Trait` requires a visible `impl Trait` for `T`. - Multiple bounds are allowed: `T: TraitA + TraitB`. - Bounds may bind associated types: `I: Iterator<Item = T>`.

Trait Coherence (Core v1)

To avoid ambiguous implementations, Core v1 applies the orphan rule: - An `impl` is valid only if either the trait or the type is defined in the current package.

Additionally: - If multiple applicable `impl` blocks could apply to the same type at a call site, the program is ill-formed. - Blanket implementations (e.g., `impl<T: Trait> OtherTrait for T`) are permitted but must not violate coherence.

Numeric Semantics (Core v1)

- Integer overflow and underflow are runtime errors and MUST trap. This applies in every build configuration; there is no mode in which overflow wraps silently.
- Division or modulo by zero is a runtime error and MUST trap.
- Floating-point operations follow IEEE-754 semantics (NaN, +/-Inf).

Type Safety Guarantees

1. **No null pointer dereferences:** Option type prevents null access
2. **No buffer overflows:** Array bounds checking
3. **No use-after-free:** Ownership system prevents dangling pointers
4. **No data races:** Borrowing rules prevent concurrent access violations
5. **No memory leaks:** Automatic memory management through ownership

Informative Future Directions (Non-Normative)

Lifetime Parameters

```
fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {  
    if x.len() > y.len() { x } else { y }  
}
```

Trait Objects

Dynamic dispatch via `dyn Trait` is a future extension; `dyn` is a reserved keyword.

Capturing Closures

Lambda expressions that capture their environment are a future extension; the `fn(...)` function types in Core v1 are non-capturing.

Conformance

A conforming Core v1 implementation MUST follow the requirements in this document. Any deviations or extensions MUST be explicitly documented by the implementation.

STARK Semantic Analysis Specification

Overview

Semantic analysis validates that programs are meaningful according to STARK's language rules. This phase occurs after parsing and before code generation.

Analysis Phases

1. Symbol Table Construction

Build symbol tables for scopes and declarations.

Scope Rules

```
// Module scope
fn module_function() { }
const MODULE_CONST: Int32 = 42;

fn example() {
  // Function scope
  let local_var = 10;

  {
    // Block scope
    let block_var = 20;
    // block_var accessible here
    // local_var accessible here
  }
  // block_var not accessible here
  // local_var still accessible
}
```

Name Resolution Order

Name resolution distinguishes *lexical* scopes (inside function bodies) from *module* scopes.

Within a function body, an unqualified name is resolved lexically: 1. Current block scope 2. Enclosing block scopes (innermost to outermost) 3. Function parameters

If not found lexically, the name is resolved at module level per the module system rules (see 07-Modules-and-Packages.md): 4. Items declared in the current module 5. Items brought into scope by use 6. Built-in names (prelude)

Unqualified names never implicitly search parent modules or the crate root; those require explicit `super::` or `crate::` paths.

Shadowing Rules

```
let x = 10;          // x: Int32
{
    let x = "hi"; // Shadows outer x, x: String
    // Inner x takes precedence
}
// Outer x visible again
```

2. Type Checking

Variable Declarations

```
let x: Int32 = 42;    // Type annotation matches literal
let y = 42;          // Type inferred as Int32
let z: String = 42;   // Error: type mismatch
```

Function Calls

```
fn add(a: Int32, b: Int32) -> Int32 { a + b }

let result = add(1, 2);    // OK: arguments match parameters
let error = add(1.0, 2);   // Error: Float64 cannot be Int32
let error2 = add(1);       // Error: wrong number of arguments
```

Assignment Compatibility

```
let mut x: Int32 = 10;
x = 20;                // OK: same type
x = "hello";           // Error: type mismatch

let y: Int32 = 10;
y = 20;                // Error: y is not mutable
```

3. Ownership and Borrowing Analysis

Move Semantics Validation

```
let s1 = String::from("hello");
let s2 = s1;           // s1 moved to s2
print(s1);             // Error: use of moved value

fn take_ownership(s: String) { }
let s3 = String::from("world");
take_ownership(s3);     // s3 moved into function
print(s3);             // Error: use of moved value
```

Borrow Checking

```
let mut s = String::from("hello");
let r1 = &s;                // Immutable borrow
let r2 = &s;                // OK: multiple immutable borrows
let r3 = &mut s;            // Error: cannot borrow mutably while
                             immutably borrowed

let mut s2 = String::from("world");
let r4 = &mut s2;           // Mutable borrow
let r5 = &s2;               // Error: cannot borrow immutably while
                             mutably borrowed
let r6 = &mut s2;           // Error: cannot have multiple mutable
                             borrows
```

Lifetime Validation

Returned references must obey the Core v1 reference rules defined in 03-Type-System.md (“References and Lifetimes”): a returned reference must derive from a reference parameter on every path, and callers treat it as having the shortest lifetime among the reference arguments it may derive from.

```
fn invalid_return() -> &Int32 {
    let x = 42;
    &x                      // Error: returning reference to local
                             variable
}

fn valid_return(x: &Int32) -> &Int32 {
    x                      // OK: returning input reference
}
```

4. Control Flow Analysis

Unreachable Code Detection

```
fn example() -> Int32 {
    return 42;
    let x = 10;           // Warning: unreachable code
}
```

Return Path Analysis

```
fn missing_return() -> Int32 {
    let x = 42;
    // Error: not all paths return a value
}

fn conditional_return(flag: Bool) -> Int32 {
    if flag {
        return 1;
    }
}
```

```

    // Error: missing return in else branch
}

fn valid_return(flag: Bool) -> Int32 {
    if flag {
        1
    } else {
        2
    }
    // OK: both branches return
}

```

Break/Continue Validation

```

fn invalid_break() {
    break;
    // Error: break outside of loop
}

fn valid_break() {
    loop {
        break;
        // OK: break inside loop
    }
}

```

5. Pattern Matching Analysis

Exhaustiveness Checking

```

enum Color { Red, Green, Blue }

fn check_color(c: Color) -> String {
    match c {
        Color::Red => "red",
        Color::Green => "green"
        // Error: non-exhaustive pattern (missing Blue)
    }
}

fn valid_match(c: Color) -> String {
    match c {
        Color::Red => "red",
        Color::Green => "green",
        Color::Blue => "blue"
        // OK: exhaustive
    }
}

```

Rules: - A match over an enum type is exhaustive if every variant is covered, or a wildcard (`_`) arm exists. - Tuple patterns are exhaustive if each element position is exhaustive for its type. - Literal patterns are exhaustive only for finite domains (e.g., `Bool` with `true` and `false`). - If a match is not exhaustive, it is a compile-time error.

Pattern Type Checking

```
let x: Int32 = 42;
match x {
  "hello" => "string",      // Error: pattern type mismatch
  42 => "number",          // OK
  _ => "other"
}
```

6. Mutability Analysis

Mutable Access Validation

Assignment to an initialized immutable variable is an error. (Exception: the single initializing assignment of a deferred-initialization `let` — see Initialization Analysis.)

```
let x = 42;
x = 43;                      // Error: x is not mutable

let mut y = 42;
y = 43;                      // OK: y is mutable

struct Point { x: Int32, y: Int32 }
let p = Point { x: 1, y: 2 };
p.x = 10;                   // Error: p is not mutable

let mut p2 = Point { x: 1, y: 2 };
p2.x = 10;                  // OK: p2 is mutable
```

7. Initialization Analysis

Use Before Initialization

```
let x: Int32;
print(x.fmt());              // Error: use of uninitialized variable

let y: Int32 = if true { 42 } else { 24 };
print(y.fmt());              // OK: y is initialized in all branches

let z: Int32;
if condition {
  z = 42;
}
print(z.fmt());              // Error: z might not be initialized
```

Rules: - `let name: Type;` declares a variable without initializing it. - A variable must be definitely assigned before any read. - All control-flow paths must assign before use. - **Deferred initialization of immutable variables:** an *uninitialized*

immutable `let` may be assigned exactly once on each control-flow path; this first assignment is initialization, not mutation. Any assignment to an already-initialized immutable variable is an error (see Mutability Analysis).

Double Initialization

```
let mut x: Int32 = 42;
x = 43;           // OK: reassignment
let x = 44;       // Error: redeclaration in same scope
```

8. Array Bounds Analysis

Static Bounds Checking

```
let arr: [Int32; 3] = [1, 2, 3];
let x = arr[2];      // OK: index in bounds
let y = arr[5];      // Error: index out of bounds (if
determinable)
```

Dynamic Bounds Checking

```
let arr: [Int32; 3] = [1, 2, 3];
let idx = get_index();
let x = arr[idx];    // Runtime bounds check required
```

9. Error Propagation Analysis

Question Mark Operator

```
fn might_fail() -> Result<Int32, String> { Ok(42) }

fn caller1() -> Result<Int32, String> {
    let x = might_fail()?;    // OK: compatible error types
    Ok(x * 2)
}

fn caller2() -> Int32 {
    let x = might_fail()?;    // Error: function doesn't return
    Result
    x * 2
}
```

Rules: - `expr?` propagates `Err` or `None` to the nearest enclosing function returning `Result<_, E>` or `Option<_>`. - The enclosing function return type must be compatible with the propagated type.

Runtime Error Semantics (Core v1)

Static analysis must identify operations whose normative runtime rule can trap and preserve the source location of that operation. Trap categories, abort behavior, and the absence of unwinding are defined solely by `CORE-V1-ABSTRACT-MACHINE.md` (`TRAP-CATEGORY-001` and `DROP-ABORT-001`). Numeric and process-specific classifications are completed by C2.9.

10. Trait Constraint Checking

Trait Bounds Validation

```
trait Display {  
    fn fmt(&self) -> String;  
}  
  
fn print_it<T: Display>(item: T) {  
    print(item.fmt());          // OK: T implements Display  
}  
  
struct Point { x: Int32, y: Int32 }  
  
print_it(Point { x: 1, y: 2 }); // Error: Point doesn't implement  
Display
```

Error Reporting

Error Message Format

```
Error: [ERROR_CODE] [BRIEF_DESCRIPTION]  
    --> file.stark:line:column  
    |  
line | source code line  
    | ^^^^ specific location  
    |  
    = help: detailed explanation  
    = note: additional information
```

Error Categories

Type Errors (E0001-E0099)

- E0001: Type mismatch
- E0002: Unknown type
- E0003: Type annotation required
- E0004: Cannot infer type
- E0005: Wrong number of arguments
- E0006: `?` operator in a function that does not return `Result` or `Option`
- E0007: Index out of bounds (determinable at compile time)

- E0008: Integer literal out of range for its type (suffix literal exceeds its suffix's representable range, or an unsuffixed literal exceeds Int64)
- E0009: Array repeat count ([value; count]) is not a compile-time constant expression

Ownership Errors (E0100-E0199)

- E0100: Use of moved value
- E0101: Borrow check failed
- E0102: Lifetime violation
- E0103: Dangling reference

Name Resolution Errors (E0200-E0299)

- E0200: Undefined variable
- E0201: Undefined function
- E0202: Undefined type
- E0203: Ambiguous name
- E0204: Duplicate definition in the same scope

Control Flow Errors (E0300-E0399)

- E0301: Missing return value
- E0302: Break outside loop
- E0303: Non-exhaustive match

(E0300 is intentionally unassigned: unreachable code is warning W0005, not an error — see “Unreachable Code Detection” above.)

Mutability and Initialization Errors (E0400-E0499)

- E0400: Assignment to immutable binding
- E0401: Use of possibly-uninitialized variable

Trait Errors (E0500-E0599)

- E0500: Trait not implemented

Warning Categories

Unused Items (W0001-W0099)

- W0001: Unused variable
- W0002: Unused function
- W0003: Unused import
- W0004: Dead code
- W0005: Unreachable code

Style Warnings (W0100-W0199)

- W0100: Non-snake_case variable

- W0101: Non-PascalCase type
- W0102: Missing documentation

Conformance

A conforming Core v1 implementation MUST follow the requirements in this document. Any deviations or extensions MUST be explicitly documented by the implementation.

STARK Core v1 Abstract Machine

Status and authority

This document is normative for Core v1. It is the sole authority for runtime evaluation, values and places, ownership transfer, temporary lifetime, destruction, references, and language traps. 03-Type-System.md defines static legality, 04-Semantic-Analysis.md defines required analyses and rejection categories, and 05-Memory-Model.md summarizes safety guarantees. None of those documents defines a competing execution model.

The abstract machine does not prescribe an interpreter frame layout, stack or heap placement, pointer representation, object layout, garbage collector, MIR shape, backend ABI, or optimizer. A conforming implementation may represent the concepts below in any way that preserves all specified behavior.

Terms and machine state

Values, objects, storage, and ownership

AM-VALUE-001. A *value* is a typed Core datum: a scalar, function reference, reference, aggregate, enum value, or value of an implementation-provided Core library type.

AM-OBJECT-001. An *object* is a value together with its identity and lifetime. Object identity is abstract. It is not a numeric address and does not change merely because ownership of the object moves between places.

AM-STORAGE-001. A *storage location* is an abstract slot capable of holding one object or being uninitialized. An *allocation* is one or more storage locations whose physical placement is not observable unless another Core rule explicitly exposes it.

AM-OWNER-001. Every live non-Copy object has exactly one owner. An owner may be a local, temporary, aggregate field or element, active enum payload, function parameter, return-transfer slot, iterator, collection, or another library value whose contract owns the object.

AM-LOCAL-001. A *local binding* names a storage location within a lexical scope. A parameter and method receiver are locals whose initialization occurs on function entry. A local is *initialized* when its location contains a live object and *moved-from* when ownership has been transferred out without replacement.

AM-TEMP-001. A *temporary* is an unnamed owner created while evaluating a full expression. Temporaries have the same exactly-once destruction obligations as named locals.

Places and projections

EXEC-PLACE-001. A *place* designates a storage location or a sublocation of an object. Core places are:

- an initialized local, parameter, or receiver;
- a field or tuple-field projection from a place;
- an array, slice, vector, map-entry, or other indexing projection whose library contract defines a place;
- a range projection denoting a slice place;
- dereference of a valid reference.

A projection is part of the place identity; resolving a place does not read or move its stored value. Field and tuple projections resolve their base first. Index projections resolve the base place, then evaluate the index. Bounds or key failure is a language trap.

An expression used where a place is required must be one of the place forms above, except that borrowing an rvalue may first materialize that rvalue in a temporary place.

Evaluation model

Results and full expressions

Evaluation of an expression produces exactly one of:

- a value;
- a place when the surrounding context requires a place;
- return, carrying a value;
- break, optionally carrying a value;
- continue;
- propagation from `?`, carrying `None` or `Err`;
- a language trap.

A *full expression* is an initializer, expression statement, return operand, break operand, condition, match scrutinee, loop iterable, function argument, aggregate field/element, assignment operand, or block tail evaluated as one enclosing operation.

EXEC-ONCE-001. Every evaluated expression and subexpression is evaluated exactly once. Static dispatch, method lookup, place resolution, auto-borrowing, and auto-dereferencing must not re-evaluate source expressions.

EXEC-DISPATCH-001. Execution invokes exactly the function, inherent method, trait method, default method, associated item, or compiler-recognized language hook selected by the normative static rules. A runtime may not substitute source-order lookup, same-named inherent lookup for a trait-qualified protocol, structural host equality/ordering, or string-name dispatch with different semantics. Dispatch selection itself has no Core-visible side effect.

Value and place contexts

MOVE-READ-001. Reading a place in a value context:

- copies the value and leaves the place initialized when its type is Copy;
- otherwise transfers ownership from the place and leaves that place moved-from;
- does not move when the construct explicitly borrows the place or a trait/operator rule specifies a reference argument.

The static rules reject any read that would move from a prohibited place, conflict with a live borrow, or read an uninitialized or moved-from place. Runtime execution of a well-typed program therefore never relies on recovering from such a read.

Evaluation order

EXEC-EVAL-001. Unless a row below states otherwise, subexpressions evaluate left to right in source order, and each completes—including its side effects, ownership transfers, normal destruction, control transfer, or trap—before the next begins.

Expression form	Normative order and result
Literal, function item, constant, unit	Produce the corresponding value. Constant evaluation is defined by 03–Type–System.md.
Local/path in value context	Resolve the named place, then apply MOVE–READ–001.
Grouping	Evaluate the grouped expression once.
Field or tuple field	Resolve the base place, append the projection, then read only if a value is required.
Integer index	Resolve the base place, evaluate the index, check bounds, append the projection, then read only if required.
Range index	Resolve the base place, evaluate the range bounds left to right, validate them, and produce a slice place/view.
Borrow &/mut	Resolve the operand place, materializing an rvalue temporary if

Expression form	Normative order and result
	permitted, then create a reference without reading the referent.
Dereference	Evaluate the reference, validate it, and produce its referent place; read only if required.
Other unary operation or cast	Evaluate the operand, then apply the operation. Numeric outcomes are completed by C2.9.
Non-short-circuit binary operation	Evaluate the left operand, then the right, then apply the operation.
<code>a && b</code>	Evaluate <code>a</code> ; evaluate <code>b</code> only when <code>a</code> is <code>true</code> .
<code>a b</code>	Evaluate <code>a</code> ; evaluate <code>b</code> only when <code>a</code> is <code>false</code> .
Range construction	Evaluate the lower bound, then upper bound, then construct the range.
Free or associated call	Resolve the callee without evaluating arguments; evaluate arguments left to right; transfer them to parameters; execute the body.
Method call	Evaluate and resolve the receiver exactly once; evaluate arguments left to right; transfer/borrow the receiver and arguments; execute the selected body.
Tuple, array, struct, enum payload	Evaluate fields/elements left to right in written order, transferring each completed value into the partial aggregate.
Array repeat [<code>value</code> ; <code>count</code>]	Evaluate <code>value</code> once; obtain the compile-time <code>count</code> ; copy the value <code>count</code> times. The static rules require <code>Copy</code> .
<code>if</code>	Evaluate the condition, then only the selected branch.
<code>match</code>	Evaluate the scrutinee exactly once; test arms in source order; evaluate only the first matching arm.
Block	Evaluate statements in order, then the optional tail expression; the block result is the tail value or <code>Unit</code> .
<code>loop</code>	Repeatedly evaluate the body until control transfers.
<code>while</code>	Evaluate the condition before each iteration and the body only when it is <code>true</code> .
<code>for</code>	

Expression form	Normative order and result
	Evaluate the iterable once, obtain its iterator, repeatedly obtain one next value and execute one iteration.
?	Evaluate the operand once; unwrap Some/Ok, or transfer None/Err toward the enclosing function.
return, break	Evaluate the optional operand before beginning the control transfer.
continue	Begin the control transfer without an operand.
Assignment	Follow EXEC-ASSIGN-001; it is the named right-before-left exception.

An expression form not admitted by the Core grammar has no Core execution semantics. A legal Core expression form omitted from this table is a specification defect, not permission for implementation-defined evaluation.

EXEC-FOR-001. Every value accepted by the static for-loop iterator rules is executed through its selected iterator protocol. The runtime may not narrow this to a privileged list of container representations. The iterable expression evaluates once, each successful step yields one iteration value, and no later step occurs after loop control exits.

Assignment and replacement

EXEC-ASSIGN-001. Simple assignment `destination = source` executes in this exact order:

1. evaluate `source` completely and retain ownership of its result;
2. resolve `destination` as a place, including evaluation and validation of its projections;
3. detach the old destination object if the place is initialized;
4. write the new object and make the destination initialized;
5. destroy the detached old object.

The new value is installed before user-observable destruction of the old value begins. If the old value's destructor traps, execution aborts with the new value already installed, although subsequent program observation is impossible because Core has no trap recovery.

Compound assignment evaluates the right operand first, resolves the destination second, reads the old destination value exactly once, applies the corresponding operation, and replaces the old value using steps 3–5 above.

If source evaluation traps, the destination is not resolved or modified. If destination resolution traps after the source has produced an owned object, the program aborts; the source object and all other live values are not destructed because DROP-ABORT-001 forbids unwinding.

```
let mut values = [0, 0];
values[next_index()] = make_value(); // make_value() completes
before next_index().
```

```
let mut values = [0];
values[1] = String::from("owned"); // RHS is produced; index
failure then aborts without Drop.
```

Aggregate construction

EXEC-AGG-001. Aggregate fields and elements become owned by a partial aggregate as each initializer completes. Declaration order determines layout-independent field identity; written initializer order determines evaluation order. Duplicate, unknown, or missing required fields are static errors and do not execute.

Tuple and positional enum payload order is index order. Named struct and enum payload initializers evaluate in written order even when that differs from declaration order. The completed object's later field-destruction order is declaration order reversed, not initializer order reversed.

EXEC-AGG-002. If a later initializer performs normal early control transfer, already completed fields are destroyed in reverse completion order before the transfer continues. If a later initializer traps, construction aborts and no completed field is destructed, because a trap never unwinds.

Core v1 has no struct-update syntax. If a later edition adds it, that edition must define its evaluation, transfer, and failure sequence rather than inheriting one implicitly.

```
let value = (Loud::new("completed"), fail()); // A trap in fail()
does not unwind completed.
```

Normal control transfer and partial evaluation

EXEC-CFLOW-001. `return`, `break`, `continue`, and `?` propagation are normal control transfers, not traps. Before a transfer leaves a scope, all live owners in that scope that are not carried by the transfer are destroyed using the normal order. The carried value is moved out before cleanup and is not destroyed by the exited scope.

For nested scopes, cleanup proceeds from the innermost exited scope outward. `break` destroys the current iteration and loop-body state before producing the loop result. `continue` destroys the current iteration and begins the next. `?` destroys scopes between the operator and the function boundary before transferring `None/Err` to the caller.

Moves, partial moves, and reinitialization

OWN-MOVE-001. Passing a non-Copy value to a by-value parameter, storing it in an aggregate, returning it, breaking with it, propagating it, or binding it by value transfers ownership. Passing or returning a Copy value copies it.

OWN-PARTIAL-001. Moving a field from an aggregate transfers only that field and marks that field moved-from. Remaining initialized fields remain individually usable and destructible. The whole aggregate cannot be read, moved, or deconstructed as fully initialized until every moved field is reinitialized. Moving a field from a type that implements `Drop` is prohibited, because its destructor requires the complete value.

Moving out of an indexed place is prohibited unless a library operation explicitly owns and removes that element (for example `Vec::remove` or `pop`). This keeps shifting/container invariants inside the owning operation.

OWN-REINIT-001. Assignment to a moved-from mutable place reinitializes that place. Assigning all moved fields reconstitutes a fully initialized aggregate. Reinitialization does not destroy the absent old value.

```
let mut text = String::from("first");
let moved = text;
text = String::from("second"); // text is initialized again; no
old text remains to drop.
```

Pattern execution

This section defines the C2.7 ownership/destruction mechanics for the pattern forms present in Core v1. C2.8 remains authoritative for pattern typing, exhaustiveness, usefulness, and any additional static restrictions.

PAT-OWN-001. A match scrutinee in value context is evaluated once. A non-Copy place scrutinee moves into a hidden scrutinee owner; a Copy scrutinee is copied. Bindings in the selected pattern take ownership of their matched subobjects unless the scrutinee itself is a reference, in which case projections remain references to the referent. The wildcard `_` takes no binding but does not suppress destruction.

PAT-DROP-001. The selected arm's bindings are arm-local owners. On every normal exit from the arm, binding locals are destroyed in reverse binding order after any arm result has moved out. Then every still-owned, unbound component of the hidden scrutinee is destroyed exactly once in reverse declaration/index order. A trap in pattern testing or the arm aborts without either cleanup.

Failed arm tests create no user-visible bindings and transfer no ownership out of the hidden scrutinee. Core v1 has no match guards or alternative `(|)` patterns.

```
struct Loud { label: String }
impl Drop for Loud { fn drop(&mut self)
{ println(self.label.as_str()); } }
let pair = (Loud { label: String::from("bound") },
           Loud { label: String::from("wildcard") });
match pair { (kept, _) => println(kept.label.as_str()), }
```

References and borrow-carrying values

Identity and derivation

REF-IDENTITY-001. A reference is a value containing:

- the identity of a live referent object;
- a projection path and, for a slice, validated bounds;
- shared or exclusive access capability;
- the static validity interval established by the type/borrow rules.

This is an abstract identity, not a promised address or pointer width. Two references alias when they designate overlapping storage of the same object.

Creating a reference does not copy or move the referent. Moving the reference value preserves its designation. Moving the referent's owner while a conflicting reference is live is rejected statically. When an ownership move is legal, object identity is preserved even if physical storage changes.

Projection, auto-borrow, and auto-dereference

REF-PROJECT-001. Dereference yields the designated place. Field, tuple, index, and slice projections through a reference extend its projection path and retain the referent's identity. Reads and writes through the resulting place observe the current object, never a snapshot.

Method auto-borrow creates a reference to the already evaluated receiver place. Auto-dereference follows references without re-evaluating the receiver. An exclusive receiver requires a mutable place and preserves write-through behavior.

Returned and receiver-derived references

REF-RETURN-001. A returned reference must designate an object whose validity, under the static shortest-input-lifetime rules, extends into the caller. Transfer across a call boundary preserves referent identity and projection. A reference derived from `&self`, `&mut self`, or another reference parameter designates the caller's object—not a parameter copy or callee-local storage—and remains valid after the callee's activation ends for the statically permitted interval.

```
struct Cell { value: Int32 }
impl Cell {
    fn value_ref(&self) -> &Int32 { &self.value }
}
let cell = Cell { value: 42 };
let reference = cell.value_ref();
```

Slice references

REF-SLICE-001. Borrowing a range-indexed place creates a view of the original base object. The view records validated start/end bounds and aliases those elements. Reads observe later legal writes; writes through an exclusive slice reference update the original object. Creating or returning a slice never clones its elements merely to satisfy reference semantics.

```
struct Numbers { items: Vec<Int32> }
impl Numbers {
    fn middle(&self) -> &[Int32] { &self.items[1..3] }
}
let numbers = Numbers { items: Vec::new() };
let middle = numbers.middle();
```

Borrow-carrying values

REF-CARRY-001. A generic aggregate or library value that contains references carries their referent identities, projections, capabilities, and validity requirements unchanged through moves, calls, returns, and pattern binding. Moving the carrier does not retarget its references. Destroying a carrier destroys owned non-reference components but does not destroy referenced objects.

Destruction

Exactly-once rule

DROP-EXACT-001. Every initialized owned object that reaches a normal destruction point is destroyed exactly once. A moved-from or never-initialized place is not destroyed. Moving a value transfers its destruction obligation to the new owner. Copy types cannot implement Drop.

For a type implementing Drop, its `Drop::drop(&mut self)` body runs first while all non-moved fields remain readable. After it returns normally, owned fields are recursively destroyed. Calling `Drop::drop` directly is prohibited; the `drop(value)` function consumes its argument and performs this sequence.

If a destructor traps, execution aborts immediately; remaining fields, siblings, locals, and outer scopes are not destructed.

Deterministic order

DROP-ORDER-001. On normal execution:

- expression-statement results are destroyed at the statement boundary;
- temporaries are destroyed in reverse creation order at the end of their full expression;
- block locals are destroyed in reverse declaration order;
- a return/break/propagated result moves out before local cleanup;
- function body locals are destroyed as their scopes exit, then parameters in reverse parameter order; a receiver is the first parameter for this purpose;

- tuple and array elements are destroyed in reverse index order;
- struct and named enum-variant fields are destroyed in reverse source declaration order;
- positional enum payloads are destroyed in reverse index order;
- only the active enum payload is destroyed;
- `Option`, `Result`, `Box`, and similar single-payload owners destroy their active payload;
- partially moved aggregates destroy only fields that remain initialized.

The order is defined in terms of source declarations and logical ownership, never a map order, address order, backend layout, or optimizer choice.

Temporaries

EXEC-TEMP-001. A temporary that is not transferred into another owner lives until the end of its full expression and is then destroyed in reverse creation order. A temporary borrowed by an unbound temporary reference lives through that statement; the static rules prevent the reference from escaping. A temporary moved into a local, aggregate, argument, return value, or other owner ceases to be a temporary and its destruction obligation transfers.

A block tail value is moved or copied out before the block's locals and remaining temporaries are destroyed. Temporaries created while evaluating a condition, match scrutinee, loop iterable, argument, or aggregate initializer follow the more specific ownership rules for that construct.

Loops and iterators

DROP-LOOP-001. A `for` iteration binding owns the yielded value for exactly one iteration. After the body exits normally, by `continue`, by `break`, by `return`, or by `?`, the binding is destroyed unless its value was moved out. Body-local destruction precedes iteration-binding destruction.

The iterator owns its unconsumed state. On normal loop completion or normal early transfer, the iterator is destroyed and therefore destroys all remaining owned elements according to its library contract. A trap aborts without this cleanup.

```
for item in values {
    if skip(item) { continue; }
    if stop(item) { break; }
    consume(item);
}
```

Collections and ownership-discarding operations

DROP-COLLECTION-001. A collection owns its stored elements and entries. Destroying or clearing a sequence/set destroys elements in reverse defined iteration order. Destroying or clearing a map destroys entries in reverse defined iteration order, destroying each value before its key. Removing an entry transfers it to the caller when the API returns it; otherwise the operation destroys it. Replacing a map value installs the new value and then destroys the old value. Rejected duplicate inputs remain owned by the operation and are destroyed.

```
let mut map: HashMap<String, Loud> = HashMap::new();
map.insert(String::from("key"), Loud::new("value"));
let removed = map.remove(&String::from("key"));
map.clear();
```

Trap termination

DROP-ABORT-001. A language trap is aborting. It performs no stack unwinding, scope cleanup, temporary cleanup, partial-aggregate cleanup, collection cleanup, or user Drop calls for live objects. Side effects and destruction completed before the trap remain observable. Core v1 has no catch or recovery mechanism.

`panic(message)` emits its specified panic diagnostic/message and then traps. The precise process exit-code mapping is owned by C2.9; the abstract exit category is trap.

```
struct Loud { label: String }
impl Drop for Loud { fn drop(&mut self)
{ println(self.label.as_str()); } }
let live = Loud { label: String::from("not dropped") };
panic("abort");
```

Trap categories

TRAP-CATEGORY-001. A *language trap* is a failure explicitly required by a normative Core rule, including explicit panic, bounds failure, invalid range boundary, division by zero, checked arithmetic failure, and failing checked conversion where the owning numeric/text rule requires a trap. It records a stable category and the source location of the operation that failed.

Allocation exhaustion, stack exhaustion, host I/O failure, OS termination, compiler limits, and target failures are not silently reclassified as language traps. C2.9 classifies those conditions under CORE-Q-017. A conforming implementation must not report an internal host panic as though it were a specified STARK trap.

Observable execution and differential comparison

OBS-COMPARE-001. For a fixed program, inputs, target contract, enabled extensions, artifact set, and external environment, execution backends are semantically equal only when all applicable observations match:

- stdout bytes in order;
- stderr bytes in order;
- abstract exit category (success, language trap, host/process failure);
- returned Core value for a harnessed function, compared by its Core value semantics;
- language-trap category;
- language-trap source identity and span;
- user-observable Drop side effects and their order;

- artifact verification result and artifact identity when verification is part of the execution request.

Wall-clock time, physical addresses, allocation strategy, stack depth, object layout, generated code bytes, and non-normative diagnostic prose are not observations. Numeric target variation, process status numbers, environmental I/O, and resource failures are compared only after C2.9 classifies them.

Conformance

A conforming Core v1 implementation must implement every rule in this document. Optimization may omit work only when it cannot change an observation in OBS-COMPARE-001, including destructor effects and trap order. The interpreter, future MIR interpreter, and native backend are implementations of this abstract machine; none is an independent semantic authority.

STARK Core v1 Memory-Safety Model

Status and authority

This chapter states Core v1 memory-safety guarantees. Static ownership, borrowing, lifetime, Copy, and Drop legality is defined by 03-Type-System.md and checked as required by 04-Semantic-Analysis.md. Runtime values, places, moves, references, temporary lifetime, destruction order, and traps are defined solely by CORE-V1-ABSTRACT-MACHINE.md.

This chapter does not define physical layout. Core v1 makes no general promise that a value is stack-allocated or heap-allocated, that a reference is a machine pointer, that a slice is two machine words, or that an enum uses a particular tag. Observable layout and target contracts are owned by C2.9.

Safety invariants

For every well-typed Core v1 program that has not terminated through a language trap or an external failure:

1. every live non-Copy object has exactly one owner;
2. a moved-from or uninitialized place cannot be read;
3. a live shared reference permits reads but not mutation through that reference;
4. a live exclusive reference excludes conflicting shared or exclusive access;
5. a reference cannot be used outside its statically valid interval;
6. dereference and projection preserve referent identity and cannot silently produce a disconnected snapshot;
7. bounds-checked places cannot access storage outside their designated object or slice;
8. every owned object that reaches normal cleanup is destroyed exactly once;
9. a value is Copy only when its complete type satisfies the Copy legality rules;

10. normal control transfer preserves ownership and cleanup obligations.

These are language guarantees, not implementation strategies. A compiler may erase borrow metadata and omit redundant runtime checks after proving that the same observable behavior and safety invariants remain.

Ownership and moves

Ownership transfer does not duplicate the transferred object. Passing or returning a non-Copy value, storing it into another owner, binding it by value, or removing it through an owning collection operation transfers the destruction obligation with the object.

Partial-move legality, prohibited indexed moves, and reinitialization checks are defined by `03-Type-System.md`. Their runtime effect is defined by abstract-machine rules `OWN-PARTIAL-001`, `OWN-REINIT-001`, and `DROP-EXACT-001`.

Borrowing and validity

Core v1 has shared and exclusive references. The static rules conservatively determine their validity without written lifetime parameters. Returned references must derive from permitted reference inputs, and borrow-carrying generic values carry the same validity constraints as their contained references.

Reference identity, projection, method receivers, returned references, slice views, moves of reference carriers, and caller/callee boundaries are defined by `REF-IDENTITY-001` through `REF-CARRY-001` in `CORE-V1-ABSTRACT-MACHINE.md`.

Destruction safety

The type system prohibits a type from implementing both Copy and Drop, prohibits explicit calls to `Drop::drop`, and prohibits partial field moves from a type whose destructor requires the complete value.

The abstract machine defines all destruction points and ordering, including locals, temporaries, parameters, fields, partial aggregates, loop bindings, iterators, collections, explicit `drop(value)`, replacement assignment, normal early transfer, and aborting traps. No physical container order or object layout may substitute for the specified logical order.

Traps and external failures

A specified language trap preserves memory safety by terminating execution; Core v1 does not unwind or run remaining destructors. Side effects completed before the trap remain observable. `TRAP-CATEGORY-001` distinguishes language traps from host panics and external failures.

Allocation exhaustion, stack exhaustion, OS termination, host I/O failure, and target limits are not memory-safety loopholes and are not automatically STARK traps. C2.9 classifies their portable guarantees.

Informative future directions (non-normative)

Written lifetime parameters, reference fields, reference counting, interior mutability, concurrency, and unsafe/native memory access are not Core v1 features. A future edition must preserve the invariants above or explicitly version its compatibility boundary.

Conformance

A conforming implementation must enforce the static rules and preserve these safety invariants while implementing the abstract machine. Deviations and extensions must be documented and tested; an implementation representation is never evidence that a different language rule applies.

STARK Standard Library Specification

Overview

The STARK standard library provides essential types, functions, and modules for core programming tasks. This specification defines the minimal standard library for the initial implementation.

Notation

The code blocks in this document are **API notation**, not compilable STARK source: function signatures are listed without bodies, and `impl` blocks mix signatures with prose comments. The *signatures and behaviors* are normative; the listing style is not. (The Core v1 grammar has no body-less function form outside `trait` blocks.)

Implementation-Provided Types

Several standard library types (`Box<T>`, `Vec<T>`, `String`, `HashMap<K, V>`) manage raw memory internally. Raw pointers and unsafe code are NOT part of the Core v1 language surface; these types are *implementation-provided*: their public APIs are

normative, their internal representation is not expressible in Core v1 source and is implementation-defined. Struct bodies shown for these types in this document are illustrative only.

Iterator and View Types

The iterator types returned by the APIs below (`VecIter<T>`, `KeysIter<K>`, `ValuesIter<V>`, `Iter<K, V>`, `Iter<T>`, `CharsIter`, `SplitIter`, `MapIter<I, U>`, `FilterIter<I>`) are implementation-provided opaque structs. Each implements `Iterator` with the obvious `Item`:

Type	Produced by	Item
<code>VecIter<T></code>	<code>Vec::iter</code>	<code>&T</code>
<code>KeysIter<K></code>	<code>HashMap::keys</code>	<code>&K</code>
<code>ValuesIter<V></code>	<code>HashMap::values</code>	<code>&V</code>
<code>Iter<K, V></code>	<code>HashMap::iter</code>	<code>(&K, &V)</code>
<code>Iter<T></code>	<code>HashSet::iter</code>	<code>&T</code>
<code>CharsIter</code>	<code>String::chars</code> , <code>str::chars</code>	<code>Char</code>
<code>SplitIter</code>	<code>String::split</code>	<code>&str</code>
<code>MapIter<I, U></code>	<code>Iterator::map</code>	<code>U</code>
<code>FilterIter<I></code>	<code>Iterator::filter</code>	<code>I::Item</code>

Iterators whose `Item` is a reference (and `CharsIter`/`SplitIter`, which borrow their source) are **borrow-carrying types** in the sense of 03-Type-System.md: they hold a live borrow of the collection they iterate and obey all reference rules (no storage in user structs, lexically scoped, cannot outlive their source).

Core Module Structure

```
std/  
├── core/           // Core language items  
├── collections/    // Data structures  
├── io/             // Input/output operations  
├── string/         // String manipulation  
├── math/           // Mathematical operations  
├── mem/            // Memory management utilities  
├── error/          // Error handling types  
└── prelude/       // Automatically imported items
```

Prelude Module

Automatically imported into every STARK program:

```
// Basic types (no import needed)  
Int8, Int16, Int32, Int64, UInt8, UInt16, UInt32, UInt64  
Float32, Float64, Bool, Char, String, str, Unit
```

```

// Essential traits
trait Copy
trait Clone
trait Drop
trait Eq
trait Ord
trait Hash
trait Default
trait Display

// Essential enums
enum Ordering {
    Less,
    Equal,
    Greater
}

enum Option<T> {
    Some(T),
    None
}

enum Result<T, E> {
    Ok(T),
    Err(E)
}

// Essential functions
fn print(value: &str)
fn println(value: &str)
fn panic(message: &str) -> !    // Never returns; see 03-Type-
System.md (Never Type)

```

Core Module (std::core)

Essential Trait Definitions

```

trait Clone {
    fn clone(&self) -> Self;
}

trait Eq {
    fn eq(&self, other: &Self) -> Bool;
}

trait Ord {
    fn cmp(&self, other: &Self) -> Ordering;
}

trait Hash {
    fn hash(&self) -> UInt64;
}

```

```

}

trait Default {
    fn default() -> Self;
}

trait Display {
    fn fmt(&self) -> String;
}

trait Drop {
    fn drop(&mut self);
}

// Copy is a marker trait (no methods); Copy: Clone.
// Soundness rules (all-fields-Copy, Copy/Drop exclusivity) are in
// 03-Type-System.md, "Copy and Drop".
trait Copy

// Num is a compiler-known marker trait implemented by exactly the
// built-in
// numeric types; it enables arithmetic operators on generic
// parameters
// (see 03-Type-System.md, "Operators and Traits"). Not user-
// implementable.
trait Num

```

Indexing Traits

The `[]` operator desugars to these traits:

```

trait Index<Idx> {
    type Output;
    fn index(&self, index: Idx) -> &Self::Output;
}

trait IndexMut<Idx> {
    fn index_mut(&mut self, index: Idx) -> &mut Self::Output;
}

```

Range Types

The range operators `..` and `..=` produce these types:

```

struct Range<T> {
    start: T,
    end: T
}

struct RangeInclusive<T> {
    start: T,
    end: T
}

```



```

}

// Ranges over integer types are iterable
impl Iterator for Range<Int32> {
    type Item = Int32;
    fn next(&mut self) -> Option<Int32>;
}
// ... similarly for the other integer types, and for
RangeInclusive

```

Memory Management

```

// Box – heap allocation (implementation-provided; internals
opaque)
struct Box<T> { /* implementation-defined */ }

impl<T> Box<T> {
    fn new(value: T) -> Box<T>;
    fn into_inner(self) -> T;
}

// Manual memory management
fn drop<T>(value: T)
fn size_of<T>() -> UInt64      // T not inferable: call as
size_of::()
fn align_of<T>() -> UInt64     // T not inferable: call as
align_of::()

```

Option Type

```

enum Option<T> {
    Some(T),
    None
}

impl<T> Option<T> {
    fn is_some(&self) -> Bool;
    fn is_none(&self) -> Bool;
    fn unwrap(self) -> T;
    fn unwrap_or(self, default: T) -> T;
    fn map<U>(self, f: fn(T) -> U) -> Option<U>;
    fn and_then<U>(self, f: fn(T) -> Option<U>) -> Option<U>;
}

```

Result Type

```

enum Result<T, E> {
    Ok(T),
    Err(E)
}

impl<T, E> Result<T, E> {

```

```

    fn is_ok(&self) -> Bool;
    fn is_err(&self) -> Bool;
    fn unwrap(self) -> T;
    fn unwrap_or(self, default: T) -> T;
    fn map<U>(self, f: fn(T) -> U) -> Result<U, E>;
    fn map_err<F>(self, f: fn(E) -> F) -> Result<T, F>;
    fn and_then<U>(self, f: fn(T) -> Result<U, E>) -> Result<U,
E>;
}

```

Collections Module (std::collections)

Vec - Dynamic Array

```

// Implementation-provided; internals opaque
struct Vec<T> { /* implementation-defined */ }

impl<T> Vec<T> {
    fn new() -> Vec<T>;
    fn with_capacity(capacity: UInt64) -> Vec<T>;
    fn push(&mut self, item: T);
    fn pop(&mut self) -> Option<T>;
    fn len(&self) -> UInt64;
    fn capacity(&self) -> UInt64;
    fn is_empty(&self) -> Bool;
    fn get(&self, index: UInt64) -> Option<&T>;
    fn get_mut(&mut self, index: UInt64) -> Option<&mut T>;
    fn insert(&mut self, index: UInt64, item: T);
    fn remove(&mut self, index: UInt64) -> T;
    fn clear(&mut self);
    fn append(&mut self, other: &mut Vec<T>);
    fn extend<I: Iterator<Item = T>>(&mut self, iter: I);
    fn iter(&self) -> VecIter<T>;
    fn as_slice(&self) -> &[T];
}

// Index access
impl<T> Index<UInt64> for Vec<T> {
    type Output = T;
    fn index(&self, index: UInt64) -> &T;
}

impl<T> IndexMut<UInt64> for Vec<T> {
    fn index_mut(&mut self, index: UInt64) -> &mut T;
}

```

HashMap<K, V> - Hash Table

```

// Implementation-provided; internals opaque
struct HashMap<K, V> { /* implementation-defined */ }

```

```
impl<K: Hash + Eq, V> HashMap<K, V> {
    fn new() -> HashMap<K, V>;
    fn with_capacity(capacity: UInt64) -> HashMap<K, V>;
    fn insert(&mut self, key: K, value: V) -> Option<V>;
    fn get(&self, key: &K) -> Option<&V>;
    fn get_mut(&mut self, key: &K) -> Option<&mut V>;
    fn remove(&mut self, key: &K) -> Option<V>;
    fn contains_key(&self, key: &K) -> Bool;
    fn len(&self) -> UInt64;
    fn is_empty(&self) -> Bool;
    fn clear(&mut self);
    fn keys(&self) -> KeysIter<K>;
    fn values(&self) -> ValuesIter<V>;
    fn iter(&self) -> Iter<K, V>;
}
```

HashSet - Hash Set

```
// Implementation-provided; internals opaque
struct HashSet<T> { /* implementation-defined */ }
```

```
impl<T: Hash + Eq> HashSet<T> {
    fn new() -> HashSet<T>;
    fn insert(&mut self, value: T) -> Bool;
    fn remove(&mut self, value: &T) -> Bool;
    fn contains(&self, value: &T) -> Bool;
    fn len(&self) -> UInt64;
    fn is_empty(&self) -> Bool;
    fn clear(&mut self);
    fn iter(&self) -> Iter<T>;
}
```

Iteration Order (Core v1)

HashMap::keys/values/iter and HashSet::iter (and any for loop over a HashMap/HashSet) MUST visit entries in **first-insertion order**: inserting a key for the first time appends it to the iteration order; inserting an already-present key updates its value without changing its position; removing a key and later re-inserting it places it at the end (as a new insertion). This requires no bound beyond Hash + Eq on the key type (unlike a key-sorted-order rule, which would additionally require K: Ord — not part of HashMap’s/HashSet’s bounds). This is normative and deterministic: two conforming implementations must produce identical iteration order for the same sequence of insertions/removals, regardless of internal storage strategy (see “Performance Notes” below, which describes storage, not iteration order).

String Module (std::string)

String Type

```
// Implementation-provided; internals opaque (UTF-8 byte buffer)
struct String { /* implementation-defined */ }

impl String {
    fn new() -> String;           // Empty string
    fn with_capacity(capacity: UInt64) -> String;
    fn from(s: &str) -> String;   // Construct from a string
    slice/literal
    fn len(&self) -> UInt64;
    fn is_empty(&self) -> Bool;
    fn push(&mut self, ch: Char);
    fn push_str(&mut self, s: &str);
    fn pop(&mut self) -> Option<Char>;
    fn clear(&mut self);
    fn chars(&self) -> CharsIter;
    fn bytes(&self) -> &[UInt8];
    fn as_str(&self) -> &str;
    fn into_bytes(self) -> Vec<UInt8>;
    fn substring(&self, start: UInt64, end: UInt64) -> &str;
    fn contains(&self, pattern: &str) -> Bool;
    fn starts_with(&self, pattern: &str) -> Bool;
    fn ends_with(&self, pattern: &str) -> Bool;
    fn find(&self, pattern: &str) -> Option<UInt64>;
    fn replace(&self, from: &str, to: &str) -> String;
    fn split(&self, delimiter: &str) -> SplitIter;
    fn trim(&self) -> &str;
    fn to_lowercase(&self) -> String;
    fn to_uppercase(&self) -> String;
}

// String literals (&str)
impl str {
    fn len(&self) -> UInt64;
    fn is_empty(&self) -> Bool;
    fn chars(&self) -> CharsIter;
    fn bytes(&self) -> &[UInt8];
    fn to_string(&self) -> String;
    // ... similar methods to String
}
```

Math Module (std::math)

Basic Operations

```
// Constants
const PI: Float64 = 3.141592653589793;
```

```

const E: Float64 = 2.718281828459045;

// Basic functions
fn abs<T: Num>(x: T) -> T
fn min<T: Ord>(a: T, b: T) -> T
fn max<T: Ord>(a: T, b: T) -> T
fn clamp<T: Ord>(value: T, min: T, max: T) -> T

// Floating point functions
fn sqrt(x: Float64) -> Float64
fn pow(base: Float64, exp: Float64) -> Float64
fn log(x: Float64) -> Float64
fn log10(x: Float64) -> Float64
fn exp(x: Float64) -> Float64

// Trigonometric functions
fn sin(x: Float64) -> Float64
fn cos(x: Float64) -> Float64
fn tan(x: Float64) -> Float64
fn asin(x: Float64) -> Float64
fn acos(x: Float64) -> Float64
fn atan(x: Float64) -> Float64
fn atan2(y: Float64, x: Float64) -> Float64

// Rounding functions
fn floor(x: Float64) -> Float64
fn ceil(x: Float64) -> Float64
fn round(x: Float64) -> Float64
fn trunc(x: Float64) -> Float64

// Random numbers (simple linear congruential generator)
struct Random {
    seed: UInt64
}

impl Random {
    fn new(seed: UInt64) -> Random;
    fn next_int(&mut self) -> UInt64;
    fn next_float(&mut self) -> Float64;
    fn range(&mut self, min: Int32, max: Int32) -> Int32;
}

```

IO Module (std::io)

Basic IO Operations

```

// Standard streams
fn print(text: &str)
fn println(text: &str)
fn eprint(text: &str)    // stderr
fn eprintln(text: &str)  // stderr

```

```

// Simple file operations
struct File { /* implementation-defined */ }

impl File {
    fn open(path: &str) -> Result<File, IOError>;
    fn create(path: &str) -> Result<File, IOError>;
    fn read_to_string(&mut self) -> Result<String, IOError>;
    fn write(&mut self, data: &[UInt8]) -> Result<UInt64,
IOError>;
    fn write_str(&mut self, text: &str) -> Result<UInt64,
IOError>;
    fn close(self) -> Result<Unit, IOError>;
}

// Error types
enum IOError {
    NotFound,
    PermissionDenied,
    AlreadyExists,
    InvalidInput,
    Other(String)
}

// Utility functions
fn read_file(path: &str) -> Result<String, IOError>
fn write_file(path: &str, content: &str) -> Result<Unit, IOError>

```

Error Module (std::error)

Error Trait

```

trait Error {
    fn message(&self) -> String;
}

// Standard error types
struct GenericError {
    message: String
}

impl Error for GenericError {
    fn message(&self) -> String {
        self.message.clone()
    }
}

```

Note: error *chaining* (a `source()` method returning a trait object) requires dyn Trait support, which is a future extension. Core v1 errors carry a message only; richer error types can embed context in their own fields.

Memory Module (std::mem)

Memory Utilities

```
// Memory operations
fn size_of<T>() -> UInt64      // Call with turbofish:
size_of::<Int32>()
fn align_of<T>() -> UInt64     // Call with turbofish:
align_of::<Int32>()
fn swap<T>(a: &mut T, b: &mut T)
fn replace<T>(dest: &mut T, src: T) -> T
fn take<T: Default>(dest: &mut T) -> T
```

Raw-pointer copy operations (`copy`, `copy_nonoverlapping`) require raw pointers and `unsafe`, which are not part of Core v1; they are deferred to a future extension.

Iterator Trait (std::iter)

Basic Iterator Interface

```
trait Iterator {
  type Item;

  fn next(&mut self) -> Option<Self::Item>;

  // Default implementations
  fn count(self) -> UInt64;
  fn collect<C: FromIterator<Self::Item>>(self) -> C;
  fn map<U>(self, f: fn(Self::Item) -> U) -> MapIter<Self, U>;
  fn filter(self, predicate: fn(&Self::Item) -> Bool) ->
FilterIter<Self>;
  fn fold<B>(self, init: B, f: fn(B, Self::Item) -> B) -> B;
  fn reduce(self, f: fn(Self::Item, Self::Item) -> Self::Item)
-> Option<Self::Item>;
  fn any(self, predicate: fn(Self::Item) -> Bool) -> Bool;
  fn all(self, predicate: fn(Self::Item) -> Bool) -> Bool;
  fn find(self, predicate: fn(&Self::Item) -> Bool) ->
Option<Self::Item>;
}

trait FromIterator<T> {
  fn from_iter<I: Iterator<Item = T>>(iter: I) -> Self;
}
```

Limitation (Core v1): the combinator parameters above are *non-capturing* function types (`fn(...)`). They accept named functions and function references but cannot capture local variables. Capturing closures are a future extension; until then, prefer explicit loops when state must be carried into the operation.

Conversion Traits

Basic Conversion

```
trait From<T> {
    fn from(value: T) -> Self;
}

trait Into<T> {
    fn into(self) -> T;
}

trait TryFrom<T> {
    type Error;
    fn try_from(value: T) -> Result<Self, Self::Error>;
}

trait TryInto<T> {
    type Error;
    fn try_into(self) -> Result<T, Self::Error>;
}

// String conversion
trait ToString {
    fn to_string(&self) -> String;
}

trait FromStr {
    type Error;
    fn from_str(s: &str) -> Result<Self, Self::Error>;
}
```

Essential Trait Implementations

Default Implementations

```
// Copy trait for basic types
impl Copy for Int8 { }
impl Copy for Int16 { }
impl Copy for Int32 { }
impl Copy for Int64 { }
impl Copy for UInt8 { }
impl Copy for UInt16 { }
impl Copy for UInt32 { }
impl Copy for UInt64 { }
impl Copy for Float32 { }
impl Copy for Float64 { }
impl Copy for Bool { }
impl Copy for Char { }
impl Copy for Unit { }
```



```
// Clone for all Copy types (blanket implementation)
impl<T: Copy> Clone for T {
    fn clone(&self) -> T { *self }
}

// Eq trait for basic types
impl Eq for Int32 {
    fn eq(&self, other: &Int32) -> Bool { *self == *other }
}
// ... similar for other types

// Ord trait for basic types
impl Ord for Int32 {
    fn cmp(&self, other: &Int32) -> Ordering {
        if *self < *other { Ordering::Less }
        else if *self > *other { Ordering::Greater }
        else { Ordering::Equal }
    }
}
}
```

Conformance Profiles

Two standard-library conformance profiles are defined. A conforming implementation **MUST** state which profile it implements.

Profile: **core-min** (MVP)

The minimum standard library for Core v1 conformance: - Prelude: primitive types, Option, Result, Ordering, essential traits (Copy, Clone, Drop, Eq, Ord, Num), print, println, panic - String and str (construction, len, push/push_str, as_str, chars) - Vec<T> (construction, push, pop, len, get/get_mut, indexing) - Range/RangeInclusive with integer Iterator impls (for for loops) - Box<T> - drop, size_of, align_of

Profile: **std-full**

Everything in this document: core-min plus HashMap, HashSet, the Iterator trait with combinators and all iterator/view types, the full String API, the math module, file IO, std::mem utilities, the conversion traits, and Hash/Default/Display/Index/IndexMut.

Items in *neither* profile (informative future work, not required for any Core v1 conformance claim): buffered IO, regular expressions, time/date, threads and concurrency primitives, Rc/Arc/RefCell.

Informative Platform Considerations (Non-Normative)

Cross-platform Abstractions

- File path handling
- Directory operations
- Environment variables
- Process spawning (future)

Behavioral Requirements (Core v1)

- Indexing `Vec<T>` with `[]` MUST perform bounds checking and MUST trap on out-of-bounds access.
 - `get/get_mut` MUST return `None` for out-of-bounds indices and MUST NOT trap.
 - Slicing an array, slice, or `Vec<T>` with a range MUST trap if the range is out of bounds or inverted.
 - `String::substring(start, end)` MUST validate UTF-8 boundaries and MUST trap on invalid boundaries or ranges.
 - IO functions MUST return `Result` with a non-Ok variant on failure and MUST NOT silently ignore errors. ## Conformance A conforming Core v1 implementation MUST implement at least the core-min profile, MUST state which profile it implements, and MUST follow the requirements in this document for every item it provides. Any deviations or extensions MUST be explicitly documented by the implementation.
-

STARK Modules and Packages Specification

Overview

This document defines the module system, visibility rules, and package resolution for the STARK core language. It is normative for Core v1.

Package Layout

A package is a directory containing a `starkpkg.json` manifest at its root.

- The manifest's `entry` field defines the root source file.
- If `entry` is omitted, the default is `src/main.stark`.

All module paths are resolved relative to the package root unless otherwise noted.

Module Declarations

Modules are declared with the `mod` keyword.

```
mod math;

mod inline {
    pub fn add(a: Int32, b: Int32) -> Int32 { a + b }
}
```

File-Based Modules

A declaration `mod name;` loads one of the following files, in order: 1. `name.stark`
2. `name/mod.stark`

The loaded file defines the contents of the module `name`.

Module Paths

Paths use `::` separators.

- `crate` refers to the package root module.
- `self` refers to the current module.
- `super` refers to the parent module.

Examples:

```
use crate::utils::math::add;
use super::config;
```

Imports (use)

The `use` statement brings names into scope.

```
use crate::utils::math::add;
use crate::utils::{math, io};
use crate::utils::math as m;
use crate::utils::*;
```

Rules: - `use` affects name resolution in the current module only. - `pub use` re-exports the imported name from the current module. - Aliases with `as` are local to the current module.

Visibility

- Items are private to their defining module by default.
- `pub` makes an item visible to parent modules and external modules.
- `priv` explicitly marks an item as private (same as default).

Visibility applies to: - Functions, structs, enums, traits, impl blocks, consts, type aliases, and modules.

Name Resolution Order

Within a module, names are resolved in the following order: 1. Local (lexical) scope — block scopes and function parameters, as defined in 04-Semantic-Analysis.md 2. Items declared in the current module 3. Items brought into scope by use 4. Built-in names (prelude)

Unqualified names do not implicitly search parent or crate scopes. Items in parent modules or the crate root require explicit `super::` or `crate::` paths. (Note: *lexical* scopes inside a function body do nest — step 1 — but *module* scopes never nest implicitly.)

Manifest Schema (starkpkg.json)

The manifest is a JSON object with the following fields:

Field	Type	Required	Rules
name	string	yes	Package name: [a-z][a-z0-9_-]*, 1–64 chars. Used as the import root for dependents.
version	string	yes	Semantic version MAJOR.MINOR.PATCH (numeric components; optional -prerelease tag).
entry	string	no	Package-root-relative path to the root source file. Default src/main.stark. MUST exist and MUST be inside the package directory.
dependencies	object	no	Map from package name to a <i>version constraint string</i> or a <i>dependency object</i> .

A dependency object has exactly one of: - { "version": "<constraint>" } — resolved from a registry or cache, or - { "path": "<relative-path>" } — a local package directory containing its own starkpkg.json.

Version constraint syntax: - "1.2.3" — exactly that version - "^1.2.3" — $\geq 1.2.3$ and $< 2.0.0$ (compatible-with) - " ≥ 1.2 , < 2.0 " — comma-separated comparator list ($\geq$, $>$, $\leq$, $<$, $=$), all of which must hold

Validation: unknown top-level fields are ignored (forward compatibility); a missing/invalid required field, a malformed constraint, a path escaping the workspace, or a dependency whose name violates the name rule is a manifest error and compilation MUST fail.

Example:

```
{
  "name": "my-app",
  "version": "0.1.0",
  "entry": "src/main.stark",
  "dependencies": {
    "tensor-lib": "^2.1.0",
    "utils": { "path": "../utils" }
  }
}
```

Packages and Dependencies

Dependencies are declared in `starkpkg.json` under `dependencies`.

Resolution order for external packages: 1. Local package source (the current package) 2. Direct dependencies from `starkpkg.json` 3. Standard library package `std`

Dependency modules are accessed via their package name:

```
use TensorLib::tensor::Tensor;
```

Dependency Version Resolution (Core v1)

- Version strings follow semantic versioning.
- If multiple versions satisfy a constraint, the highest version MUST be selected.
- If no version satisfies a constraint, compilation MUST fail with an error.
- The source of packages (registry, cache, or local path) is implementation-defined, but the chosen version MUST be reported in build output.

Standard Library

The standard library is available under the `std` package name:

```
use std::io;
use std::collections::Vec;
```

Cycles

- **Package dependency cycles** — direct ($A \rightarrow A$) or transitive ($A \rightarrow B \rightarrow C \rightarrow A$) — are a compile-time error.

- **Module declarations** (`mod`) form a tree rooted at the package entry file, so `mod` cycles cannot occur; declaring the same module twice is an error.
- **use imports** between modules of the same package MAY be mutually recursive (module A may use items from B and vice versa); this is not a cycle error.

Errors

- Importing an unknown module or item is a compile-time error.
- Accessing a private item outside its module is a compile-time error.
- Ambiguous imports are a compile-time error unless aliased. ## Conformance A conforming Core v1 implementation MUST follow the requirements in this document. Any deviations or extensions MUST be explicitly documented by the implementation.